

BookReview_EDA_Satyajit_Ghana

December 11, 2020

Book Review - An Exploratory Data Analysis and Data Transformation Notebook

Author: Satyajit Ghana

```
[1]: import gdown

url = 'https://drive.google.com/uc?id=1UPZiTughL3iDtPwreoUs_SX-LfVktrI3'
output = 'BX-CSV-Dump.zip'
gdown.download(url, output, quiet=False)
```

Downloading...

From: https://drive.google.com/uc?id=1UPZiTughL3iDtPwreoUs_SX-LfVktrI3

To: /content/BX-CSV-Dump.zip

26.1MB [00:01, 23.9MB/s]

```
[1]: 'BX-CSV-Dump.zip'
```

```
[2]: ! unzip BX-CSV-Dump.zip
```

Archive: BX-CSV-Dump.zip

inflating: BX-Book-Ratings.csv

inflating: BX-Books.csv

inflating: BX-Users.csv

1 Book Crossing EDA

```
[67]: %matplotlib inline

import scipy
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

```

from sklearn.ensemble import IsolationForest
from sklearn.impute import SimpleImputer

sns.set()
palette = sns.color_palette("icefire")

plt.style.use('ggplot')

sns.set_context("talk")

```

1.1 Loading, Cleaning and Merging the Dataset

Reading the csv files

The dataset is semi-colon separated instead of semi-colon separated, and als the incoding is ISO-8859-1

```

[4]: users = pd.read_csv(
    '/content/BX-Users.csv',
    names=['user_id', 'location', 'age'],
    sep=';',
    skiprows=1,
    encoding='ISO-8859-1',
    low_memory=False,
    error_bad_lines=False
)
users

```

```

[4]:
   user_id      location  age
0         1  nyc, new york, usa  NaN
1         2  stockton, california, usa  18.0
2         3  moscow, yukon territory, russia  NaN
3         4  porto, v.n.gaia, portugal  17.0
4         5  farnborough, hants, united kingdom  NaN
...
278853  278854  portland, oregon, usa  NaN
278854  278855  tacoma, washington, united kingdom  50.0
278855  278856  brampton, ontario, canada  NaN
278856  278857  knoxville, tennessee, usa  NaN
278857  278858  dublin, n/a, ireland  NaN

```

[278858 rows x 3 columns]

parse the datatypes properly

```

[5]: users.dtypes

```

```
[5]: user_id      int64
      location    object
      age         float64
      dtype: object
```

A quick look at the numeric attributes of the users table

```
[6]: users.describe().T
```

```
[6]:
```

	count	mean	std	...	50%	75%
max						
user_id	278858.0	139429.500000	80499.515020	...	139429.5	209143.75
278858.0						
age	168096.0	34.751434	14.428097	...	32.0	44.00
244.0						

```
[2 rows x 8 columns]
```

Data Inconsistency

We notice that the max age is 244, and min age is 0, age cannot be 244 ! so let's fix that, also minimum age cannot be 0, this is likely a mistake during data collection, and missing age values were probably just replaced by 0, so we'll have to fix that

one way is to simply replace the inconsistent values with the mean of the data

```
[212]: users.loc[(users.age > 100) | (users.age < 5), 'age'] = np.nan
        users.age = users.age.fillna(users.age.mean())
```

```
[8]: users['age'] = users['age'].astype(np.uint8)
```

```
[9]: users['age'].describe()
```

```
[9]: count    278858.000000
      mean      34.446733
      std       10.551712
      min        5.000000
      25%       29.000000
      50%       34.000000
      75%       35.000000
      max      100.000000
      Name: age, dtype: float64
```

checking for any NA values

```
[10]: users.isna().sum()
```

```
[10]: user_id    0
      location    0
```

```
age          0
dtype: int64
```

Now we'll read the books data, same way as before

```
[11]: books = pd.read_csv(
        '/content/BX-Books.csv',
        names=['isbn', 'book_title', 'book_author', 'year_of_publication', 'publisher', 'img_s', 'img_m', 'img_l'],
        sep=';',
        skiprows=1,
        encoding='ISO-8859-1',
        low_memory=False,
        error_bad_lines=False
    )
books
```

```
[11]:
```

	isbn	...	img_1
0	0195153448	...	http://images.amazon.com/images/P/0195153448.0...
1	0002005018	...	http://images.amazon.com/images/P/0002005018.0...
2	0060973129	...	http://images.amazon.com/images/P/0060973129.0...
3	0374157065	...	http://images.amazon.com/images/P/0374157065.0...
4	0393045218	...	http://images.amazon.com/images/P/0393045218.0...
...
271374	0440400988	...	http://images.amazon.com/images/P/0440400988.0...
271375	0525447644	...	http://images.amazon.com/images/P/0525447644.0...
271376	006008667X	...	http://images.amazon.com/images/P/006008667X.0...
271377	0192126040	...	http://images.amazon.com/images/P/0192126040.0...
271378	0767409752	...	http://images.amazon.com/images/P/0767409752.0...

[271379 rows x 8 columns]

parse the data types properly

```
[12]: books.dtypes
```

```
[12]: isbn          object
book_title        object
book_author        object
year_of_publication object
publisher          object
img_s             object
img_m             object
img_l             object
dtype: object
```

Dropping Unnecessary Values

drop ['img_s', 'img_m', 'img_l'] since they are not useful for us

```
[13]: books = books.drop(['img_s', 'img_m', 'img_l'], axis=1)
```

year_of_publication should be a integer

```
[14]: books['year_of_publication'] = pd.to_numeric(books['year_of_publication'],  
↳errors='coerce')
```

```
[15]: books.loc[(books['year_of_publication'] == 0) | (books['year_of_publication'] >  
↳2008), 'year_of_publication'] = np.nan  
books['year_of_publication'] = books['year_of_publication'].  
↳fillna(round(books['year_of_publication'].mean()))  
books['year_of_publication'] = pd.to_numeric(books['year_of_publication'],  
↳downcast='unsigned')
```

Checking for any NA values

```
[16]: books.isna().sum()
```

```
[16]: isbn                0  
book_title              0  
book_author             1  
year_of_publication     0  
publisher               2  
dtype: int64
```

Since the NA rows are very few, we'll simply drop them

```
[17]: books = books.dropna()
```

```
[18]: books.describe().T
```

```
[18]:
```

	count	mean	std	...	50%	75%
max						
year_of_publication	271376.0	1993.692427	8.248715	...	1995.0	2000.0
2008.0						

```
[1 rows x 8 columns]
```

Read the ratings dataset as usual

```
[19]: ratings = pd.read_csv(  
    '/content/BX-Book-Ratings.csv',  
    names=['user_id', 'isbn', 'book_rating'],  
    sep=';',  
    skiprows=1,  
    encoding='ISO-8859-1',  
    low_memory=False,  
    error_bad_lines=False
```

```
)  
ratings
```

```
[19]:
```

	user_id	isbn	book_rating
0	276725	034545104X	0
1	276726	0155061224	5
2	276727	0446520802	0
3	276729	052165615X	3
4	276729	0521795028	6
...
1149775	276704	1563526298	9
1149776	276706	0679447156	0
1149777	276709	0515107662	10
1149778	276721	0590442449	10
1149779	276723	05162443314	8

```
[1149780 rows x 3 columns]
```

```
[20]: ratings['book_rating'] = ratings['book_rating'].astype(np.uint8)
```

```
[21]: ratings.dtypes
```

```
[21]: user_id      int64  
      isbn        object  
      book_rating  uint8  
      dtype: object
```

check for any NA values

```
[22]: ratings.isna().sum()
```

```
[22]: user_id      0  
      isbn        0  
      book_rating  0  
      dtype: int64
```

There are some inconsistent values, which we will fix later

```
[23]: ratings.describe().T.astype(np.int32)
```

```
[23]:
```

	count	mean	std	min	25%	50%	75%	max
user_id	1149780	140386	80562	2	70345	141010	211028	278854
book_rating	1149780	2	3	0	0	0	7	10

Join the three datasets based on user_id and isbn as the key

```
[24]: temp = pd.merge(users, ratings, on='user_id')  
      temp = pd.merge(temp, books, on='isbn')
```

```
dataset = temp.copy()
```

```
[25]: dataset
```

```
[25]:
```

	user_id	...	publisher
0	2	...	Oxford University Press
1	8	...	HarperFlamingo Canada
2	11400	...	HarperFlamingo Canada
3	11676	...	HarperFlamingo Canada
4	41385	...	HarperFlamingo Canada
...
1031167	278851	...	Simon & Schuster
1031168	278851	...	Broadway Books
1031169	278851	...	Lone Star Books
1031170	278851	...	Kqed Books
1031171	278851	...	American Map Corporation

```
[1031172 rows x 9 columns]
```

Split the location into city, state and country and replacing missing location details with just n/a

We might use these attributes at a later stage of the assignment in collaborative filtering, for now it seems useful so we'll keep them.

```
[26]: location = dataset['location'].str.split(', ', n=2, expand=True)
location.columns = ['city', 'state', 'country']
location = location.fillna('n/a')
```

```
[27]: dataset['city'] = location['city'] ; dataset['state'] = location['state'] ;
      ↪ dataset['country'] = location['country']
```

```
[28]: dataset = dataset.drop(['location'], axis=1)
```

```
[29]: dataset.describe().T.astype(np.int32)
```

```
[29]:
```

	count	mean	std	...	50%	75%	max
user_id	1031172	140594	80524	...	141210	211426	278854
age	1031172	36	10	...	34	41	100
book_rating	1031172	2	3	...	0	7	10
year_of_publication	1031172	1995	7	...	1997	2001	2008

```
[4 rows x 8 columns]
```

checking for NA values

```
[30]: dataset.isna().sum()
```

```
[30]: user_id      0
      age         0
      isbn        0
      book_rating 0
      book_title   0
      book_author  0
      year_of_publication 0
      publisher    0
      city         0
      state        0
      country      0
      dtype: int64
```

So we have 1,031,172 values in total, that's a lot !

```
[31]: dataset.shape
```

```
[31]: (1031172, 11)
```

```
[32]: dataset.dtypes
```

```
[32]: user_id      int64
      age         uint8
      isbn        object
      book_rating  uint8
      book_title   object
      book_author  object
      year_of_publication uint16
      publisher    object
      city         object
      state        object
      country      object
      dtype: object
```

This will be the final dataset we will be working with !

```
[33]: dataset.head(5)
```

```
[33]:   user_id  age  isbn ...  city  state country
0      2    18 0195153448 ... stockton  california    usa
1      8    34 0002005018 ... timmins   ontario  canada
2   11400   49 0002005018 ... ottawa    ontario  canada
3   11676   34 0002005018 ...    n/a      n/a    n/a
4   41385   34 0002005018 ... sudbury    ontario  canada
```

[5 rows x 11 columns]

```
[34]: dataset.info()
```



```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1031172 entries, 0 to 1031171
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   user_id               1031172 non-null  int64
1   age                   1031172 non-null  uint8
2   isbn                  1031172 non-null  object
3   book_rating           1031172 non-null  uint8
4   book_title            1031172 non-null  object
5   book_author           1031172 non-null  object
6   year_of_publication   1031172 non-null  uint16
7   publisher             1031172 non-null  object
8   city                  1031172 non-null  object
9   state                 1031172 non-null  object
10  country               1031172 non-null  object
dtypes: int64(1), object(7), uint16(1), uint8(2)
memory usage: 74.7+ MB

```

```
[35]: cleaned_data = dataset.copy()
```

```
[204]: # dataset = cleaned_data.copy()
```

1.2 Analyzing the Feature Space

```

[213]: def get_skewness(data, columns):
        """returns the skewness and kurtosis of the specified attributes"""
        skewness = data[columns].skew()
        kurtosis = data[columns].kurtosis()

        df = {
            'skewness': skewness.values,
            'kurtosis': kurtosis.values
        }

        dataframe = pd.DataFrame(data=df, index=columns)

        return dataframe


```

```

[206]: f, axes = plt.subplots(ncols = 4, figsize=(25, 5))
sns.kdeplot(x="book_rating", data=dataset, ax=axes[0]).set_title('rating_distr')
sns.kdeplot(x="book_rating", data=dataset[dataset['book_rating'] != 0],
    ↪ax=axes[1]).set_title('explicit rating_distr')
sns.boxplot(y="book_rating", data=dataset, orient='v', ax=axes[2]).
    ↪set_title('rating_distr')

```

```
sns.boxplot(y="book_rating", data=dataset[dataset['book_rating'] != 0],
            orient='v', ax=axes[3]).set_title('explicit rating_distr')
plt.show()
```



Something we can clearly notice in `rating_distr` is that 0 ratings are a lot ! and when we drop them to get `explicit_rating_distr`, the distribution is lot better, even the box plot is better, so it makes sense to remove the 0 rated values

Why would someone rate a book 0 ? this likely are the values for user that forgot to rate the book, people are lazy right ?

For normally distributed data, the skewness should be about zero. For unimodal continuous distributions, a skewness value greater than zero means that there is more weight in the right tail of the distribution. The function `skewtest` can be used to determine if the skewness value is close enough to zero, statistically speaking.

```
[207]: get_skewness(dataset, ['age', 'book_rating'])
```

```
[207]:          skewness  kurtosis
age          0.832497  1.343472
book_rating  0.752445 -1.214994
```

```
[208]: dataset[['age', 'book_rating']].describe().T
```

```
[208]:          count      mean      std  min  25%  50%  75%  max
age      1031172.0  36.232554  10.413914  5.0  31.0  34.0  41.0  100.0
book_rating  1031172.0   2.839005   3.854142  0.0   0.0   0.0   7.0   10.0
```

We can remove 0 ratings, since these are unrated, and why would someone rate a book as 0 ?

```
[209]: dataset = dataset[dataset['book_rating'] != 0]
```

We notice that skewness of the `book_rating` changes from a strong positive, to negative !

```
[210]: get_skewness(dataset, ['age', 'book_rating'])
```

```
[210]:          skewness  kurtosis
age          0.859061  1.644862
```

```
book_rating -0.661283  0.120288
```

```
[211]: dataset[['age', 'book_rating']].describe().T
```

```
[211]:
```


	count	mean	std	min	25%	50%	75%	max
age	383849.0	35.856535	10.363322	5.0	31.0	34.0	40.0	100.0
book_rating	383849.0	7.626702	1.841335	1.0	7.0	8.0	9.0	10.0

1.2.1 Univariate Analysis of Attributes


Here we plot the Distribution, Box Plot, Violin Plot and QQ Plot for the numeric attributes `age` and `book_rating`

```
[133]: def plot_univariate(dataset, column_name, supitle = None, kde_only = True):
        f, axes = plt.subplots(ncols = 4, figsize=(25, 5))
        if kde_only:
            sns.kdeplot(x=column_name, data=dataset, ax=axes[0]).
            ↪set_title(f'{column_name}_distr')
        else:
            sns.histplot(x=column_name, data=dataset, kde=True, ax=axes[0]).
            ↪set_title(f'{column_name}_distr')
            sns.boxplot(y=column_name, data=dataset, orient='v', ax=axes[1]).
            ↪set_title(f'{column_name}_box_plot')
            sns.violinplot(y=column_name, data=dataset, orient='v', ax=axes[2]).
            ↪set_title(f'{column_name}_violin_plot')
            scipy.stats.probplot(dataset[column_name], dist="norm", plot=axes[3])
        if supitle:
            plt.suptitle(supitle)
            plt.subplots_adjust(top=0.80)
        plt.show()
```

```
[134]: plot_univariate(dataset=dataset, column_name='book_rating', supitle='Book_
        ↪Rating')
```



```
[135]: plot_univariate(dataset=dataset, column_name='age', subtitle='Age')
```



Inference

- **book_rating**

We see that the distribution of books is not normal, also the box plot shows outliers in the lower region, ratings from 1 to 4, the probability plot also confirms this, this is something we deal with outlier analysis later

- **age**

Age has a peak density about 30-34, the box-plot shows a lot of outliers, and the qq plot also

1.3 Outlier Analysis and removing them

We'll try to infer the outliers using the `IsolationForest` algorithm,

The `IsolationForest` 'isolates' observations by randomly selecting a feature and then randomly selecting a split value between the maximum and minimum values of the selected feature.

Since recursive partitioning can be represented by a tree structure, the number of splittings required to isolate a sample is equivalent to the path length from the root node to the terminating node.

This path length, averaged over a forest of such random trees, is a measure of normality and our decision function.

Random partitioning produces noticeably shorter paths for anomalies. Hence, when a forest of random trees collectively produce shorter path lengths for particular samples, they are highly likely to be anomalies. (Scikit-learn)

1. It classify the data point to outlier and not outliers and works great with very high dimensional data.
2. It works based on decision tree and it isolate the outliers.
3. If the result is -1, it means that this specific data point is an outlier. If the result is 1, then it means that the data point is not an outlier.

```
[136]: def plot_isolation_forest(data, columns, subtitle = None):
        ncols = len(columns)
        fig, axes = plt.subplots(nrows=1, ncols=ncols, figsize=(len(columns) * 8, 5))
        isolation_forest = IsolationForest(contamination='auto')
        for i, column in enumerate(columns):
            isolation_forest.fit(data[column].values.reshape(-1, 1))


            xx = np.linspace(data[column].min(), data[column].max(), len(data[column])).reshape(-1, 1)
            anomaly_score = isolation_forest.decision_function(xx)
            outlier = isolation_forest.predict(xx)

            axes[i].plot(xx, anomaly_score, label='anomaly_score')
            axes[i].fill_between(xx.T[0], np.min(anomaly_score), np.max(anomaly_score),
                                where=outlier == -1, color='r',
                                alpha=.4, label='outlier_region')

            axes[i].legend()
            axes[i].set_title(column)

        if subtitle:
            plt.suptitle(subtitle)
            plt.subplots_adjust(top=0.80)

[137]: plot_isolation_forest(data=dataset, columns=['age', 'book_rating'],
                               subtitle='Isolation Forest')
```



Isolation Forest shows us that a lot of data in our range contains anomalies

1.3.1 Drop Outliers with IQR

In this method by using Inter Quartile Range (IQR), we detect outliers. IQR tells us the variation in the data set. Any value, which is beyond the range of $-1.5 \times IQR$ to $1.5 \times IQR$ treated as outliers

```
[138]: def drop_outliers(data, columns):  
        data_new = data.copy()  
  
        for column in columns:  
            iqr = 1.5 * (np.percentile(data_new[column], 75) - np.  
→percentile(data_new[column], 25))  
            data_new.drop(data_new[data_new[column] > (iqr + np.  
→percentile(data_new[column], 75))].index, inplace=True)  
            data_new.drop(data_new[data_new[column] < (np.  
→percentile(data_new[column], 25) - iqr)].index, inplace=True)  
  
        return data_new
```

```
[139]: dataset_wo_outliers = drop_outliers(dataset, ['age', 'book_rating'])
```

```
[140]: dataset_wo_outliers[['age', 'book_rating']].describe().T
```

```
[140]:
```


	count	mean	std	min	25%	50%	75%	max
age	343213.0	34.534607	7.697323	17.0	31.0	34.0	37.0	53.0
book_rating	343213.0	7.742178	1.668787	4.0	7.0	8.0	9.0	10.0

```
[141]: get_skewness(dataset_wo_outliers, ['age', 'book_rating'])
```


```
[141]:
```

	skewness	kurtosis
age	0.356242	0.169093
book_rating	-0.349030	-0.808479


```
[142]: plot_univariate(dataset=dataset_wo_outliers, column_name='book_rating',  
→suptitle='Book Rating')
```



```
[143]: plot_univariate(dataset=dataset_wo_outliers, column_name='age', subtitle='Age')
```



```
[144]: plot_isolation_forest(data=dataset_wo_outliers, columns=['age', 'book_rating'],
    ↳ subtitle='Isolation Forest')
```



1.3.2 BoxCox

A Box Cox transformation is a transformation of a non-normal dependent variables into a normal shape. Normality is an important assumption for many statistical techniques; if your data isn't normal, applying a Box-Cox means that you are able to run a broader number of tests.

The BoxCox transformation is given by

$$y = \begin{cases} (x^{\lambda} - 1) / \lambda, & \text{for } \lambda > 0 \\ \log(x), & \text{for } \lambda = 0 \end{cases}$$

The confidence limits returned when alpha is provided give the interval where:

$$lf(\hat{\lambda}) - lf(\lambda) < \frac{1}{2}\chi^2(1 - \alpha, 1)$$

with llf the log-likelihood function and the chi-squared function.

```
[145]: def apply_boxcox(data, columns):  
        data_new = data.copy()  
  
        for column in columns:  
            data_new[column], _ = scipy.stats.boxcox(data_new[column].astype(np.  
→float32), lambda=None)  
  
        return data_new
```

```
[146]: dataset_boxcox = apply_boxcox(dataset, ['age', 'book_rating'])
```

```
[147]: dataset_boxcox[['age', 'book_rating']].describe().T
```

```
[147]:
```

	count	mean	std	...	50%	75%	max
age	383849.0	7.096895	1.012567	...	7.009771	7.591456	11.576612
book_rating	383849.0	22.629185	9.147474	...	23.642124	29.426302	35.763058

[2 rows x 8 columns]

```
[148]: get_skewness(dataset_boxcox, ['age', 'book_rating'])
```


```
[148]:
```

	skewness	kurtosis
age	0.042891	1.143758
book_rating	-0.170664	-0.837900


```
[149]: plot_univariate(dataset=dataset_boxcox, column_name='book_rating',  
→suptitle='Book Rating')
```



```
[150]: plot_univariate(dataset=dataset_boxcox, column_name='age', suptitle='Age')
```

```
[151]: plot_isolation_forest(data=dataset_boxcox, columns=['age', 'book_rating'],
    ↳suptitle='Isolation Forest')
```



1.3.3 Imputation

Another types of outliers are caused by missing values, for which one type of imputation algorithm is univariate, which imputes values in the i -th feature dimension using only non-missing values in that feature dimension (e.g. `impute.SimpleImputer`). (Scikit-learn)

```
[152]: def apply_imputation(data, columns):
    data_new = data.copy()

    imputer = SimpleImputer(missing_values=np.nan, strategy='mean')

    for column in columns:
        data_new[column] = imputer.fit_transform(data_new[column].astype(np.
    ↳float32).values.reshape(-1, 1))

    return data_new
```

```
[153]: dataset_imputed = apply_imputation(data=dataset, columns=['age', 'book_rating'])
```

```
[154]: dataset_imputed[['age', 'book_rating']].describe().T
```

```
[154]:
```


	count	mean	std	min	25%	50%	75%	max
age	383849.0	35.856533	10.366258	5.0	31.0	34.0	40.0	100.0
book_rating	383849.0	7.626702	1.838844	1.0	7.0	8.0	9.0	10.0

```
[155]: get_skewness(dataset_imputed, ['age', 'book_rating'])
```


```
[155]:
```

	skewness	kurtosis
age	0.859061	1.644862
book_rating	-0.661283	0.120288


```
[156]: plot_univariate(dataset=dataset_imputed, column_name='book_rating',
    ↳suptitle='Book Rating')
```



```
[157]: plot_univariate(dataset=dataset_imputed, column_name='age', suptitle='Age')
```



```
[158]: plot_isolation_forest(data=dataset_imputed, columns=['age', 'book_rating'],
    ↳suptitle='Isolation Forest')
```



You'll observe that imputation didn't help us much, since our data does not have `nan` values, but this is a demonstration that imputation can also be used

1.4 Data Transformation

```
[223]: dataset_trans = dataset_wo_outliers.copy()
```

```
[259]: def apply_function(data, columns, function):
        data_new = data.copy()

        for column in columns:
            data_new[column] = function(data_new[column])

        return data_new
```

1.4.1 Min-Max Normalization

This estimator scales and translates each feature individually such that it is in the given range on the training set, e.g. between zero and one.

The transformation is given by:

$$X_{std} = (X - X.min(axis=0)) / (X.max(axis=0) - X.min(axis=0))$$

$$X_{scaled} = X_{std} * (max - min) + min$$

where `min`, `max` = `feature_range`.

This transformation is often used as an alternative to zero mean, unit variance scaling.

The mathematical formula being:

$$x_{scaled} = \frac{x - x_{min}}{x_{max} - x_{min}}$$

```
[266]: def min_max_scaling(series):
        scaler = MinMaxScaler()
        scaled_vals = scaler.fit_transform(series.values.reshape(-1, 1)).reshape(-1)
        return scaled_vals
```

```
[267]: scaled_dset = apply_function(dataset_trans, ['age', 'book_rating'],
        ↪min_max_scaling)
scaled_dset[['age', 'book_rating']]
```

```
[267]:
```

	age	book_rating
1	0.472222	0.166667
3	0.472222	0.666667
5	0.361111	0.666667
8	0.472222	0.833333
9	0.472222	0.833333
...
1031166	0.444444	0.500000
1031168	0.444444	0.166667
1031169	0.444444	0.500000
1031170	0.444444	0.500000
1031171	0.444444	1.000000

[343213 rows x 2 columns]

```
[268]: scaled_dset[['age', 'book_rating']].describe().T
```

```
[268]:
```

	count	mean	std	...	50%	75%	max
age	343213.0	0.487072	0.213815	...	0.472222	0.555556	1.0
book_rating	343213.0	0.623696	0.278131	...	0.666667	0.833333	1.0


[2 rows x 8 columns]

```
[227]: get_skewness(scaled_dset, ['book_rating', 'age'])
```


```
[227]:
```

	skewness	kurtosis
book_rating	-0.349030	-0.808479
age	0.356242	0.169093

```
[229]: plot_univariate(dataset=scaled_dset, column_name='book_rating',
        ↪suptitle='MinMax Scaler')
```



```
[230]: plot_univariate(dataset=scaled_dset, column_name='age', subtitle='MinMax_Scaler')
```



Inference

- The skewness has now becomes -0.34 for `book_rating` and 0.35 for `age`

1.4.2 Z-Score Standardization

Standardize features by removing the mean and scaling to unit variance

The standard score of a sample x is calculated as:

$$z = \frac{x - \mu}{\sigma}$$

where μ is the mean of the training samples or zero if `with_mean=False`, and σ is the standard deviation of the training samples or one if `with_std=False`.

Centering and scaling happen independently on each feature by computing the relevant statistics on the samples in the training set. Mean and standard deviation are then stored to be used on later data using `transform`.

Standardization of a dataset is a common requirement for many machine learning estimators: they might behave badly if the individual features do not more or less look like standard normally distributed data (e.g. Gaussian with 0 mean and unit variance).

For instance many elements used in the objective function of a learning algorithm (such as the RBF kernel of Support Vector Machines or the L1 and L2 regularizers of linear models) assume that all

features are centered around 0 and have variance in the same order. If a feature has a variance that is orders of magnitude larger than others, it might dominate the objective function and make the estimator unable to learn from other features correctly as expected. (Scikit-learn)

```
[269]: def z_score_scaling(series):
        scaler = StandardScaler()
        scaled_vals = scaler.fit_transform(series.values.reshape(-1, 1)).reshape(-1)
        return scaled_vals
```

```
[270]: scaled_dset = apply_function(dataset_trans, ['age', 'book_rating'],
        ↪z_score_scaling)
scaled_dset[['age', 'book_rating']]
```

```
[270]:          age  book_rating
1      -0.069454   -1.643219
3      -0.069454    0.154497
5      -0.589116    0.154497
8      -0.069454    0.753736
9      -0.069454    0.753736
...
1031166  -0.199369   -0.444741
1031168  -0.199369   -1.643219
1031169  -0.199369   -0.444741
1031170  -0.199369   -0.444741
1031171  -0.199369    1.352974
```

[343213 rows x 2 columns]

```
[271]: scaled_dset[['age', 'book_rating']].describe().T
```


```
[271]:          count          mean         std  ...          50%          75%          max
age          343213.0 -1.057090e-13  1.000001  ... -0.069454  0.320293  2.398940
book_rating  343213.0  9.450267e-16  1.000001  ...  0.154497  0.753736  1.352974
```

[2 rows x 8 columns]


```
[234]: get_skewness(scaled_dset, ['book_rating', 'age'])
```

```
[234]:          skewness  kurtosis
book_rating -0.349030 -0.808479
age          0.356242  0.169093
```

```
[235]: plot_univariate(dataset=scaled_dset, column_name='book_rating',
        ↪suptitle='Z-Score Scaler')
```



```
[236]: plot_univariate(dataset=scaled_dset, column_name='age', subtitle='Z-Score_
↳Scaler')
```



1.4.3 Decimal Scaling

Decimal scaling is a data normalization technique like Z score, Min-Max, and normalization with standard deviation. Decimal scaling is a data normalization technique. In this technique, we move the decimal point of values of the attribute. This movement of decimal points totally depends on the maximum value among all values in the attribute.

The formula is given by:

$$v'_i = \frac{v_i}{10^j}$$

```
[272]: def decimal_scaling(series):
        p = series.max()
        q = len(str(abs(p)))
        scaled_vals = series.values / 10 ** q

        return scaled_vals
```

```
[273]: scaled_dset = apply_function(dataset_trans, ['age', 'book_rating'],
↳decimal_scaling)
scaled_dset[['age', 'book_rating']]
```

```
[273]:
```

	age	book_rating
1	0.34	0.05
3	0.34	0.08
5	0.30	0.08
8	0.34	0.09
9	0.34	0.09
...
1031166	0.33	0.07
1031168	0.33	0.05
1031169	0.33	0.07
1031170	0.33	0.07
1031171	0.33	0.10

[343213 rows x 2 columns]

```
[274]: scaled_dset[['age', 'book_rating']].describe().T
```

```
[274]:
```


	count	mean	std	min	25%	50%	75%	max
age	343213.0	0.345346	0.076973	0.17	0.31	0.34	0.37	0.53
book_rating	343213.0	0.077422	0.016688	0.04	0.07	0.08	0.09	0.10

```
[240]: get_skewness(scaled_dset, ['book_rating', 'age'])
```


```
[240]:
```

	skewness	kurtosis
book_rating	-0.349030	-0.808479
age	0.356242	0.169093

```
[241]: plot_univariate(dataset=scaled_dset, column_name='book_rating',
    ↳suptitle='Decimal Scaler')
```



```
[242]: plot_univariate(dataset=scaled_dset, column_name='age', suptitle='Decimal
    ↳Scaler')
```

1.5 Data Normality

When a metric variable fails to satisfy the assumption of normality, homogeneity of variance, or linearity, we may be able to correct the deficiency by using a transformation

- If the distribution of a variable is negatively skewed the adjustment of the values reverses, or reflects, the distribution so that it becomes positively skewed. The transformations are then computed on the values in the positively skewed distribution.
- Reflection is computed by subtracting all of the values for a variable from one plus the absolute value of maximum value for the variable. This results in a positively skewed distribution with all values larger than zero.

Which transformation?

The main criterion in choosing a transformation is: what works with the data? As examples above indicate, it is important to consider as well two questions.

What makes physical (biological, economic, whatever) sense, for example in terms of limiting behaviour as values get very small or very large? This question often leads to the use of logarithms.

Can we keep dimensions and units simple and convenient? If possible, we prefer measurement scales that are easy to think about. The cube root of a volume and the square root of an area both have the dimensions of length, so far from complicating matters, such transformations may simplify them. Reciprocals usually have simple units, as mentioned earlier. Often, however, somewhat complicated units are a sacrifice that has to be made.

It is not always necessary or desirable to transform a data set to resemble a normal distribution. However, if symmetry or normality are desired, they can often be induced through one of the power transformations.;

1.5.1 Natural Log Transform

Log transformation or $\log_e x$ is a strong transformation with a major effect on distribution shape, it is commonly used for reducing **right skewness** and is often appropriate for measured values. It cannot be applied to zero or negative values. One unit on a logarithm scale means a multiplication by the base of logarithms being used. Exponential growth or decline.

$$x_{new} = \ln x$$

```
[275]: def natural_log_transform(series):
        trans_vals = np.log(series.astype(np.float32))

        return trans_vals
```

```
[276]: trans_dset = apply_function(dataset_trans, ['age', 'book_rating'],
        ↪natural_log_transform)
trans_dset[['age', 'book_rating']]
```

```
[276]:
```

	age	book_rating
1	3.526361	1.609438
3	3.526361	2.079442
5	3.401197	2.079442
8	3.526361	2.197225
9	3.526361	2.197225
...
1031166	3.496508	1.945910
1031168	3.496508	1.609438
1031169	3.496508	1.945910
1031170	3.496508	1.945910
1031171	3.496508	2.302585

[343213 rows x 2 columns]

```
[277]: trans_dset[['age', 'book_rating']].describe().T
```

```
[277]:
```

	count	mean	std	...	50%	75%	max
age	343213.0	3.511600	0.229122	...	3.526361	3.610918	3.970292
book_rating	343213.0	2.021872	0.234512	...	2.079442	2.197225	2.302585


[2 rows x 8 columns]

```
[247]: get_skewness(trans_dset, ['book_rating', 'age'])
```


```
[247]:
```

	skewness	kurtosis
book_rating	-0.747712	-0.236441
age	-0.405089	0.544089

```
[248]: plot_univariate(dataset=trans_dset, column_name='book_rating',
        ↪suptitle='Natural Log Transform')
```



```
[249]: plot_univariate(dataset=trans_dset, column_name='age', subtitle='Natural Log Transform')
        ↪ Transform')
```



1.5.2 Square Root Transform

The square root, x to $x^{(1/2)} = \sqrt{x}$, is a transformation with a moderate effect on distribution shape: it is weaker than the logarithm and the cube root. It is also used for reducing right skewness, and also has the advantage that it can be applied to zero values. Note that the square root of an area has the units of a length. It is commonly applied to counted data, especially if the values are mostly rather small.

$$x_{new} = \sqrt{x}$$

```
[278]: def sqrt_transform(series):
        trans_vals = np.sqrt(series.astype(np.float32))

        return trans_vals
```

```
[279]: trans_dset = apply_function(dataset_trans, ['age', 'book_rating'],
        ↪ sqrt_transform)
        trans_dset[['age', 'book_rating']]
```

```
[279]:      age  book_rating
1      5.830952      2.236068
```

```

3          5.830952      2.828427
5          5.477226      2.828427
8          5.830952      3.000000
9          5.830952      3.000000
...
1031166    5.744563      2.645751
1031168    5.744563      2.236068
1031169    5.744563      2.645751
1031170    5.744563      2.645751
1031171    5.744563      3.162278

```

[343213 rows x 2 columns]

```
[280]: trans_dset[['age', 'book_rating']].describe().T
```

```

[280]:          count      mean      std  ...      50%      75%      max
age          343213.0  5.846551  0.656963  ...  5.830952  6.082763  7.280110
book_rating  343213.0  2.765608  0.310675  ...  2.828427  3.000000  3.162278

```

[2 rows x 8 columns]

```
[253]: get_skewness(trans_dset, ['book_rating', 'age'])
```


```

[253]:          skewness  kurtosis
book_rating -0.542093 -0.576979
age         -0.008925  0.214740


```

NOTE: Skewness has a value of ~0 for age using Square Root Transform

```
[254]: plot_univariate(dataset=trans_dset, column_name='book_rating', subtitle='Square Root Transform')
```



```
[255]: plot_univariate(dataset=trans_dset, column_name='age', subtitle='Square Root Transform')
```



1.5.3 Inverse Square Root Transformation

The inverse square root, $1/\sqrt{x}$ is a very strong transformation with a drastic effect on distribution shape. It can not be applied to zero values or negative values, it is not useful unless all values are positive.

$$x_{new} = \frac{1}{\sqrt{x}}$$

```
[281]: def inv_sqrt_transform(series):
        trans_vals = np.power(series.astype(np.float32), -1/2)

        return trans_vals
```

```
[285]: trans_dset = apply_function(dataset_trans, ['age', 'book_rating'],
        ↪inv_sqrt_transform)
trans_dset[['age', 'book_rating']]
```

```
[285]:
```

	age	book_rating
1	0.171499	0.447214
3	0.171499	0.353553
5	0.182574	0.353553
8	0.171499	0.333333
9	0.171499	0.333333
...
1031166	0.174078	0.377964
1031168	0.174078	0.447214
1031169	0.174078	0.377964
1031170	0.174078	0.377964
1031171	0.174078	0.316228

[343213 rows x 2 columns]

```
[286]: trans_dset[['age', 'book_rating']].describe().T
```

```
[286]:
```

	count	mean	std	...	50%	75%	max
age	343213.0	0.173505	0.020453	...	0.171499	0.179605	0.242536
book_rating	343213.0	0.366663	0.044961	...	0.353553	0.377964	0.500000


[2 rows x 8 columns]

```
[288]: get_skewness(trans_dset, ['book_rating', 'age'])
```


```
[288]:
```

	skewness	kurtosis
book_rating	0.966190	0.242251
age	0.824243	1.218674

```
[289]: plot_univariate(dataset=trans_dset, column_name='book_rating',
↳suptitle='Inverse Square Root Transform')
```



```
[290]: plot_univariate(dataset=trans_dset, column_name='age', suptitle='Inverse Square_
↳Root Transform')
```



1.6 Exploratory Data Analysis

A lot of implicit EDA has already been done above, for univariate variables, now we will go on with trying to make some meaning of other attributes of the dataset, such as `book_title`, `book_author` and `year_of_publication`

```
[ ]: # to plot values in barplot, https://stackoverflow.com/a/56780852
def show_values_on_bars(axes, h_v="v", space=0.4):
    def _show_on_single_plot(ax):
        if h_v == "v":
            for p in ax.patches:
                _x = p.get_x() + p.get_width() / 2
                _y = p.get_y() + p.get_height()
                value = int(p.get_height())
                ax.text(_x, _y, value, ha="center")
        elif h_v == "h":
            for p in ax.patches:
                _x = p.get_x() + p.get_width() + float(space)
                _y = p.get_y() + p.get_height()
                value = int(p.get_width())
                ax.text(_x, _y, value, ha="left")


    if isinstance(axes, np.ndarray):
        for idx, ax in np.ndenumerate(axes):
            _show_on_single_plot(ax)
    else:
        _show_on_single_plot(axes)
```

1.6.1 Top 25 Years of Publication

here we try to explore what were the years with huge amount of book publication, this is done by simply counting the number of book in a given year

```
[ ]: eda = dataset['year_of_publication'].copy().value_counts().head(25).
    ↪reset_index()
eda.columns = ['year_of_publication', 'count']
eda = eda.sort_values(by=['count'], ascending=False)
```

```
[ ]: plt.figure(figsize=(7, 10))
splot = sns.barplot(x='count', y='year_of_publication', data=eda,
    ↪order=eda['year_of_publication'], orient='h', palette=palette)
show_values_on_bars(splot, h_v="h")
plt.title('Top 25 Years of Publication')
plt.show()
```




Turns out 2002 was the when highest number of books were published, i.e. 37985

1.6.2 Top 10 Reviewed Books

Here we try to find the Books which were reviewed the most !


```
[ ]: eda = dataset['book_title'].value_counts().head(10).reset_index()
eda.columns = ['book_title', 'count']
```

```
[ ]: plt.figure(figsize=(7, 10))
plot = sns.barplot(x='count', y='book_title', data=eda,
    ↳order=eda['book_title'], orient='h', palette=palette)
show_values_on_bars(plot, h_v="h")
plt.title('Top 10 Reviewed Books')
plt.show()
```



The Lovely Bone: A Novel turns out to be most reviewed !

1.6.3 Top 25 Avg. Rated Books

If i want to buy something, one thing i try is to sort the values based on average rating, we try to do similar things for our books dataset, we find the highest reviewed books and then calculate their average rating and sort them again

```
[ ]: rating_count = dataset['book_title'].value_counts().head(25).reset_index().
    ↳sort_values(by='book_title').reset_index(drop=True)
rating_count.columns = ['book_title', 'rating_count']
rating_count.head(5)
```

```
[ ]:
      book_title  rating_count
0  The Girls' Guide to Hunting and Fishing      259
1                      The Testament      261
2                      Timeline      263
3      The Catcher in the Rye      265
4      To Kill a Mockingbird      267
```

```
[ ]: rating_sum = dataset[dataset['book_title'].isin(rating_count['book_title'])].
      ↳groupby(['book_title'])['book_rating'].sum().reset_index().
      ↳sort_values(by=['book_title'])
rating_sum.columns = ['book_title', 'rating_sum']
rating_sum.head(5)
```

```
[ ]:
      book_title  rating_sum
0      A Painted House      2708.0
1      Angels & Demons      2485.0
2      Bridget Jones's Diary      2875.0
3  Divine Secrets of the Ya-Ya Sisterhood: A Novel      2544.0
4      Girl with a Pearl Earring      2219.0
```


```
[ ]: avg_rating = pd.merge(rating_count, rating_sum, on='book_title')
avg_rating['rating_avg'] = avg_rating['rating_sum'] / avg_rating['rating_count']
avg_rating = avg_rating.sort_values(by='rating_avg', ascending=False).
      ↳reset_index(drop=True)
```

```
[ ]: avg_rating.head(5)
```

```
[ ]:
      book_title  ...  rating_avg
0  Harry Potter and the Prisoner of Azkaban (Book 3)  ...      9.043321
1      To Kill a Mockingbird  ...      8.977528
2  Harry Potter and the Sorcerer's Stone (Harry P...  ...      8.936508
3  Harry Potter and the Chamber of Secrets (Book 2)  ...      8.840491
4      The Secret Life of Bees  ...      8.477833
```

[5 rows x 4 columns]

```
[ ]: plt.figure(figsize=(7, 10))
      splot = sns.barplot(x='rating_avg', y='book_title', data=avg_rating,
      ↳order=avg_rating['book_title'], orient='h', palette=palette)
      show_values_on_bars(splot, h_v="h")
      plt.title('Top 25 Avg. Rated Books')
      plt.show()
```



And it turns out Harry Potter and the Prisoner of Azkaban (Book 3) is average Rated 9 !


1.6.4 Top 10 Reviewed Authors

When buying books, author are a major selling point as well, a famous author has more likelihood to be bought again if he publishes a new book

```
[ ]: author_count = dataset['book_author'].value_counts().head(10).reset_index()
author_count.columns = ['book_author', 'rating_count']
author_count.head(5)
```

```
[ ]:
   book_author  rating_count
0   Stephen King         4639
1   Nora Roberts         2938
2    John Grisham         2550
3  James Patterson         2387
4    J. K. Rowling         1746
```

```
[ ]: plt.figure(figsize=(7, 10))
splot = sns.barplot(x='rating_count', y='book_author', data=author_count,
                    order=author_count['book_author'], orient='h', palette=palette)
show_values_on_bars(splot, h_v="h")
plt.title('Top 10 Reviewed Authors')
plt.show()
```



Stephen King books are the most reviewed !

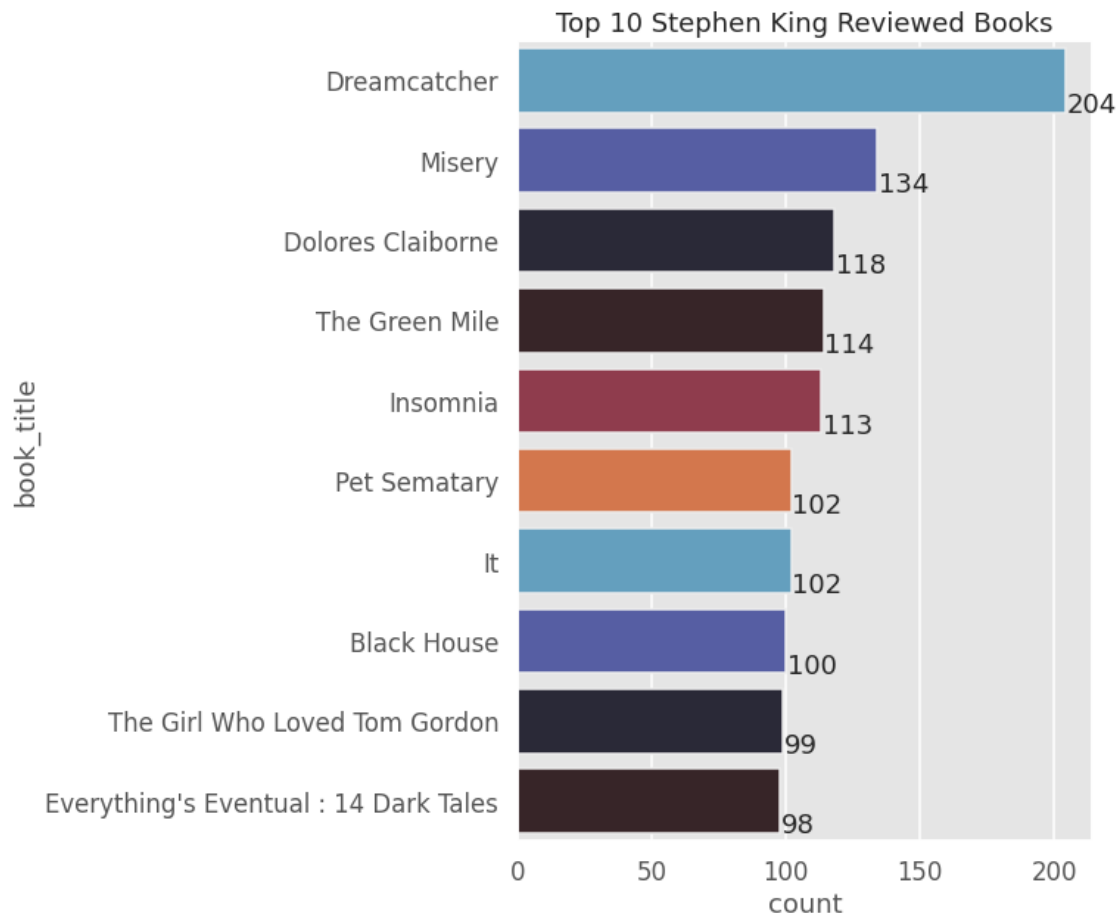
1.6.5 Analysing Stephen King Books

Since Stephen King is one of the popular authors, we can dig more about it, starting with his most reviewed books !

```
[ ]: stephen_king = dataset[dataset['book_author'] == 'Stephen King']
```

```
[ ]: eda = stephen_king['book_title'].value_counts().head(10).reset_index()
eda.columns = ['book_title', 'count']
```

```
[ ]: plt.figure(figsize=(7, 10))
      sns.barplot(x='count', y='book_title', data=eda,
                  order=eda['book_title'], orient='h', palette=palette)
      show_values_on_bars(splot, h_v="h")
      plt.title('Top 10 Stephen King Reviewed Books')
      plt.show()
```



It turns out Stephen King's book **Dreamcatcher** is the most reviewed book of his ! which was reviewed 204 times !

Now we can do the same for his books as we did before, i.e. calculate the average ratings of his books and sort them to see which book of Stephen King was the most average rated

```
[ ]: rating_count = stephen_king['book_title'].value_counts().head(10).reset_index().
      sort_values(by='book_title').reset_index(drop=True)
      rating_count.columns = ['book_title', 'rating_count']
      rating_sum = stephen_king[stephen_king['book_title'].
                               isin(rating_count['book_title'])].groupby(['book_title'])['book_rating'].
      sum().reset_index().sort_values(by='book_title')
```

```

rating_sum.columns = ['book_title', 'rating_sum']
avg_rating = pd.merge(rating_count, rating_sum, on='book_title')
avg_rating['rating_avg'] = avg_rating['rating_sum'] / avg_rating['rating_count']
avg_rating = avg_rating.sort_values(by='rating_avg', ascending=False).
    ↳reset_index(drop=True)

```

```
[ ]: avg_rating.head(5)
```

```


[ ]:
   book_title  rating_count  rating_sum  rating_avg
0  The Green Mile         114        972.0      8.526316
1             It          102        829.0      8.127451
2        Misery          134       1082.0      8.074627
3  Pet Sematary          102        808.0      7.921569
4      Insomnia          113        895.0      7.920354

```

```

[ ]: plt.figure(figsize=(7, 10))
      splot = sns.barplot(x='rating_avg', y='book_title', data=avg_rating,
    ↳order=avg_rating['book_title'], orient='h', palette=palette)
      show_values_on_bars(splot, h_v="h")
      plt.title('Top 10 Avg. Rated Stephen King Books')
      plt.show()

```



Green Mile was Stephen King's top avg. rated book with a avg. rating of 8

2 Bibliography

1. Scikit-learn: Machine Learning in Python, Pedregosa et al., JMLR 12, pp. 2825-2830, 2011.
2. G.E.P. Box and D.R. Cox, "An Analysis of Transformations", Journal of the Royal Statistical Society B, 26, 211-252 (1964).
3. <http://fmwww.bc.edu/repec/bocode/t/transint.html>

3 Convert this Notebook to PDF

```
[214]: %%capture
! apt update
! apt install texlive-xetex texlive-fonts-recommended
↪ texlive-generic-recommended
```

```
[215]: import subprocess
import shlex
```

Convert to PDF

```
[ ]: s = subprocess.Popen(shlex.split(
    f'jupyter nbconvert /content/BookReview_EDA_Satyajit_Ghana.ipynb --to pdf'
), shell = False, stdout = subprocess.PIPE, stderr = subprocess.PIPE)
s.wait()
s.stdout.read()
```

Convert to L^AT_EX

```
[219]: s = subprocess.Popen(shlex.split(
    f'jupyter nbconvert /content/BookReview_EDA_Satyajit_Ghana.ipynb --to latex'
), shell = False, stdout = subprocess.PIPE, stderr = subprocess.PIPE)
s.wait()
s.stdout.read()
```

```
[219]: b''
```

```
[220]: ! zip -r BookReview_EDA_Satyajit_Ghana.zip BookReview_EDA_Satyajit_Ghana_files_
↪BookReview_EDA_Satyajit_Ghana.tex
```

```
    adding: BookReview_EDA_Satyajit_Ghana_files/ (stored 0%)
    adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_95_0.png
(deflated 4%)
    adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_77_0.png
(deflated 7%)
    adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_195_0.png
(deflated 13%)
    adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_126_0.png
(deflated 6%)
    adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_104_0.png
(deflated 7%)
    adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_114_0.png
(deflated 4%)
    adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_162_0.png
(deflated 6%)
    adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_182_0.png
```


(deflated 14%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_155_0.png
(deflated 5%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_169_0.png
(deflated 6%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_105_0.png
(deflated 4%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_63_0.png
(deflated 6%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_134_0.png
(deflated 5%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_83_0.png
(deflated 4%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_177_0.png
(deflated 13%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_76_0.png
(deflated 5%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_206_0.png
(deflated 14%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_113_0.png
(deflated 7%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_94_0.png
(deflated 7%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_93_0.png
(deflated 5%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_141_0.png
(deflated 6%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_201_0.png
(deflated 13%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_190_0.png
(deflated 12%)
adding:
BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_112_0.png

(deflated 5%)

adding:

BookReview_EDA_Satyajit_Ghana_files/BookReview_EDA_Satyajit_Ghana_103_0.png

(deflated 5%)

adding: BookReview_EDA_Satyajit_Ghana.tex (deflated 43%)

[]: