

# BookReview-Classification-Unsupervised

January 20, 2021

```
[!]: ! pip install --upgrade scikit-learn
[!]: ! pip install scikit-learn-extra
[!]: ! pip install hdbscan
[!]: ! python -m spacy download en_core_web_sm
```

---

## 1 Book Crossing - Classification (Unsupervised)

```
[3]: %matplotlib inline

import scipy
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

sns.set()
palette = sns.color_palette("icefire")

plt.style.use('ggplot')

sns.set_context("talk")
```

```
[127]: dataset = pd.read_csv('book_crossing.classification.cleaned.csv')
```

```
[128]: dataset['age'] = dataset['age'].astype(np.float64)
dataset['book_rating'] = dataset['book_rating'].astype('category')
dataset['book_title'] = dataset['book_title'].astype('category')
dataset['book_author'] = dataset['book_author'].astype('category')
dataset['year_of_publication'] = dataset['year_of_publication'].astype(np.
    ↪float64)
dataset['publisher'] = dataset['publisher'].astype('category')
dataset['country'] = dataset['country'].astype('category')
```

```
[136]: dataset['book_title'].cat.categories.shape
```

```
[136]: (132033,)
```

```
[139]: dataset['book_author'].cat.categories.shape
```

```
[139]: (60652,)
```

```
[140]: dataset['publisher'].cat.categories.shape
```

```
[140]: (11311,)
```

```
[141]: dataset['country'].cat.categories.shape
```

```
[141]: (51,)
```

```
[6]: dataset.head()
```

```
[6]:   age book_rating ... publisher country
0  34.0         mid ... HarperFlamingo Canada  canada
1  30.0         high ... HarperFlamingo Canada  canada
2  34.0         high ... HarperFlamingo Canada  canada
3  34.0         high ... HarperFlamingo Canada  canada
4  34.0         high ... HarperFlamingo Canada  canada
```

```
[5 rows x 7 columns]
```

```
[7]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 364570 entries, 0 to 364569
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	age	364570 non-null	float64
1	book_rating	364570 non-null	category
2	book_title	364570 non-null	category
3	book_author	364570 non-null	category
4	year_of_publication	364570 non-null	float64
5	publisher	364570 non-null	category
6	country	364570 non-null	category

```
dtypes: category(5), float64(2)
```

```
memory usage: 19.1 MB
```

```
[8]: c_dataset = dataset["book_title"].astype(str) + " " + \
      dataset["book_author"].astype(str) + " " + \
      dataset["publisher"].astype(str) + " " + \
      dataset["year_of_publication"].astype(str) + " " + \
      dataset["age"].astype(str) + " " + \
      dataset["country"].astype(str)
```

```
[9]: for ex in c_dataset.sample(frac=0.2)[:5]:
      print(ex)
```

The Vanishing Vampire (The Accidental Monsters , No 1) David Lubar Scholastic 1997.0 12.0 usa  
 REMEMBER ME Mary Higgins Clark Simon & Schuster 1994.0 34.0 usa  
 The Mummy or Ramses the Damned Anne Rice Ballantine Books 1991.0 22.0 usa  
 Bittersweet Rain Sandra Brown Warner Books 2000.0 34.0 usa  
 Arthur Stephen R. Lawhead Zondervan Publishing Company 1996.0 21.0 usa

```
[92]: small_dataset = dataset.copy().sample(frac=0.03)
```

```
[93]: small_dataset.shape
```

```
[93]: (10937, 7)
```

```
[94]: c_dataset_small = small_dataset["book_title"].astype(str) + " " + \
      small_dataset["book_author"].astype(str) + " " + \
      small_dataset["publisher"].astype(str) + " " + \
      small_dataset["year_of_publication"].astype(str) + " " + \
      small_dataset["age"].astype(str) + " " + \
      small_dataset["country"].astype(str)
```

```
[95]: import spacy
      from sklearn.base import BaseEstimator, TransformerMixin
      import en_core_web_sm

      nlp = en_core_web_sm.load()
```

```
[96]: from sklearn import set_config
      from sklearn.compose import ColumnTransformer
      from sklearn.pipeline import Pipeline
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import classification_report, plot_confusion_matrix, \
      →confusion_matrix, accuracy_score, balanced_accuracy_score
      from sklearn.cluster import KMeans, DBSCAN, Birch, MiniBatchKMeans, \
      →SpectralClustering, AgglomerativeClustering, MeanShift, AffinityPropagation, \
      →OPTICS
      from sklearn.decomposition import TruncatedSVD

      set_config(display='diagram')
```

```
[97]: class SpacyVectorTransformer(BaseEstimator, TransformerMixin):
      def __init__(self, nlp):
          self.nlp = nlp
          self.dim = 300

      def fit(self, X, y):
          return self

      def transform(self, X):
          # Doc.vector defaults to an average of the token vectors.
          # https://spacy.io/api/doc#vector
```

```
return [self.nlp(text).vector for text in X]
```

```
[98]: X, y = c_dataset, dataset['book_rating']
```

```
[99]: X_small, y_small = c_dataset_small, small_dataset['book_rating']
```

```
[100]: target_names = ['low', 'mid', 'high']
```

## 1.1 Unsupervised Models

- KMedoids
- AgglomerativeClustering
- DBSCAN
- KMeansClustering
- HDBSCAN
- MiniBatchKMeans

```
[102]: """  
km = KMeans(n_clusters=3) # works (needs remap of output)  
dbscan = DBSCAN()  
birch = Birch() # crash  
mbkm = MiniBatchKMeans() # works (needs remap of output)  
sc = SpectralClustering() # crash  
ac = AgglomerativeClustering() # sparse not supported  
ms = MeanShift(bandwidth=3) # sparse not supported  
ap = AffinityPropagation() # crash  
oo = OPTICS() # sparse not supported  
"""
```

```
[102]: '\nkm = KMeans(n_clusters=3) # works (needs remap of output)\nndbscan =  
DBSCAN()\nbirch = Birch() # crash\nmbkm = MiniBatchKMeans() # works (needs remap  
of output)\nsc = SpectralClustering() # crash\nnac = AgglomerativeClustering() #  
sparse not supported\nms = MeanShift(bandwidth=3) # sparse not supported\nnap =  
AffinityPropagation() # crash\nnoo = OPTICS() # sparse not supported\n'
```

```
[103]: from time import time  
  
def fit_model(algorithm, data):  
  
    t1 = time()  
  
    X, y = data  
  
    print(f'\nStarted Training {algorithm.__class__.__name__} on X: {X.shape} y:  
→ {y.shape}')  
  
    embeddings_pipeline = Pipeline(  
        steps=[  
            ("mean_embeddings", SpacyVectorTransformer(nlp=nlp)),
```

```

        ("reduce_dim", TruncatedSVD(50)),
        ("clusterer", algorithm),
    ]
)

# train the model
embeddings_pipeline.fit(X, y)

print(f"\nEvaluating model on X_test: {X.shape} y_test: {y.shape}")

# test the model
y_true = y.copy()

if isinstance(embeddings_pipeline['clusterer'], (AgglomerativeClustering,
→DBSCAN, OPTICS, HDBSCAN)):
    y_pred = embeddings_pipeline['clusterer'].labels_
else:
    y_pred = embeddings_pipeline.predict(X)

y_pred = np.array(list(map(lambda x: "low" if x == 0 else "mid" if x == 1
→else "high", y_pred)))

# get the classification report
print(f"\nClassification Report for {algorithm.__class__.__name__}")
print(classification_report(y_true, y_pred, target_names=target_names,
→labels=target_names))

acc_score = accuracy_score(y_true, y_pred)
bal_score = balanced_accuracy_score(y_true, y_pred)

print(f"\nAccuracy Score: {acc_score}")
print(f"Balanced Accuracy Score: {bal_score}")

print()
# show the confusion matrix
cmmat_table = pd.DataFrame({'y_true': y_true, 'y_pred': y_pred})
conmat = pd.crosstab(cmmat_table.y_true, cmmat_table.y_pred,
→rownames=['Actual'], colnames=['Predicted'], margins=True, normalize='all')
ax = plt.axes()
sns.set(rc={'figure.figsize': (9, 7)})
sns.heatmap(conmat, annot=True, ax=ax)
ax.set_title(f'{algorithm.__class__.__name__}')
plt.show()
print()

t2 = time()

```

```

print(f'Trained {algorithm.__class__.__name__} in {(t2 - t1)}s')

return embeddings_pipeline

```

## 1.2 KMedoids Clustering

```

[107]: from sklearn_extra.cluster import KMedoids
[108]: kmedoids = KMedoids(n_clusters=3, max_iter=1)
[109]: clf = fit_model(algorithm=kmedoids, data=(X_small, y_small))

```

Started Training KMedoids on X: (10937,) y: (10937,)


Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for KMedoids

	precision	recall	f1-score	support
low	0.02	0.34	0.04	231
mid	0.38	0.30	0.34	4333
high	0.60	0.31	0.41	6373
accuracy			0.31	10937
macro avg	0.33	0.32	0.26	10937
weighted avg	0.50	0.31	0.37	10937

Accuracy Score: 0.30785407332906645

Balanced Accuracy Score: 0.3184037656596693



Trained KMedoids in 223.7921495437622s

```
[110]: clf
```

```
[110]: Pipeline(steps=[('mean_embeddings',
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at
0x7ff253d3b6d8>)),
                      ('reduce_dim', TruncatedSVD(n_components=50)),
                      ('clusterer', KMedoids(max_iter=1, n_clusters=3))])
```

### 1.3 Agglomerative Clustering

```
[111]: aggcl = AgglomerativeClustering(n_clusters=3)
```

```
[112]: clf = fit_model(algorithm=aggcl, data=(X_small, y_small))
```

Started Training AgglomerativeClustering on X: (10937,) y: (10937,)


Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for AgglomerativeClustering

	precision	recall	f1-score	support
low	0.02	0.37	0.04	231
mid	0.37	0.31	0.34	4333
high	0.57	0.24	0.34	6373
accuracy			0.27	10937
macro avg	0.32	0.31	0.24	10937
weighted avg	0.48	0.27	0.33	10937

Accuracy Score: 0.2729267623662796

Balanced Accuracy Score: 0.30877414843520845





Trained AgglomerativeClustering in 119.64335346221924s

```
[113]: clf
```

```
[113]: Pipeline(steps=[('mean_embeddings',  
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at  
                        0x7ff253d3b6d8>)),  
                      ('reduce_dim', TruncatedSVD(n_components=50)),  
                      ('clusterer', AgglomerativeClustering(n_clusters=3))])
```

## 1.4 DBSCAN

```
[114]: dbscan = DBSCAN(n_jobs=-1)
```

```
[115]: clf = fit_model(algorithm=dbscan, data=(X_small, y_small))
```

Started Training DBSCAN on X: (10937,) y: (10937,)


Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for DBSCAN

	precision	recall	f1-score	support
low	0.00	0.00	0.00	231
mid	0.20	0.00	0.00	4333
high	0.58	1.00	0.74	6373
accuracy			0.58	10937
macro avg	0.26	0.33	0.25	10937
weighted avg	0.42	0.58	0.43	10937

Accuracy Score: 0.5818780287098839

Balanced Accuracy Score: 0.3328872224199264



Trained DBSCAN in 116.67007970809937s

```
[116]: clf
```

```
[116]: Pipeline(steps=[('mean_embeddings',
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at
0x7ff253d3b6d8>)),
                      ('reduce_dim', TruncatedSVD(n_components=50)),
                      ('clusterer', DBSCAN(n_jobs=-1))])
```

## 1.5 KMeans Clustering

```
[117]: kmeans = KMeans(n_clusters=3, n_init=3)
```

```
[118]: clf = fit_model(algorithm=kmeans, data=(X_small, y_small))
```

Started Training KMeans on X: (10937,) y: (10937,)


Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for KMeans

	precision	recall	f1-score	support
low	0.02	0.40	0.04	231
mid	0.39	0.26	0.32	4333
high	0.60	0.35	0.44	6373
accuracy			0.32	10937
macro avg	0.34	0.34	0.27	10937
weighted avg	0.50	0.32	0.38	10937

Accuracy Score: 0.31535155892840816

Balanced Accuracy Score: 0.33658550914710594



Trained KMeans in 222.76584196090698s

```
[119]: clf
```

```
[119]: Pipeline(steps=[('mean_embeddings',  
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at  
                        0x7ff253d3b6d8>)),  
                        ('reduce_dim', TruncatedSVD(n_components=50)),  
                        ('clusterer', KMeans(n_clusters=3, n_init=3))])
```

## 1.6 HDBSCAN

```
[120]: from hdbscan import HDBSCAN
```

```
[121]: hscan = HDBSCAN(min_cluster_size=3)
```

```
[122]: clf = fit_model(algorithm=hscan, data=(X_small, y_small))
```

Started Training HDBSCAN on X: (10937,) y: (10937,)

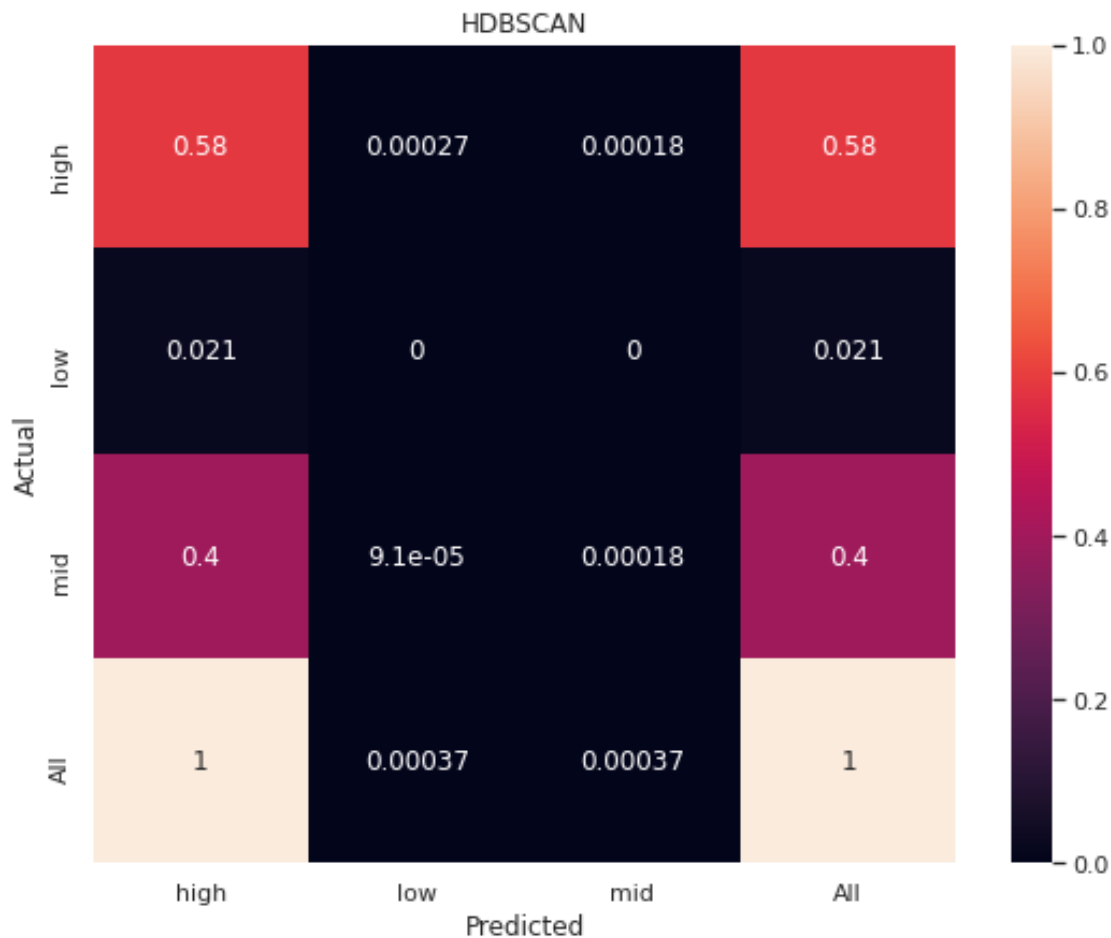
Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for HDBSCAN

	precision	recall	f1-score	support
low	0.00	0.00	0.00	231
mid	0.50	0.00	0.00	4333
high	0.58	1.00	0.74	6373
accuracy			0.58	10937
macro avg	0.36	0.33	0.25	10937
weighted avg	0.54	0.58	0.43	10937

Accuracy Score: 0.5824266252171528

Balanced Accuracy Score: 0.3332256713684369



Trained HDBSCAN in 121.39391756057739s

```
[123]: clf
```

```
[123]: Pipeline(steps=[('mean_embeddings',
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at
0x7ff253d3b6d8>)),
                        ('reduce_dim', TruncatedSVD(n_components=50)),
                        ('clusterer', HDBSCAN(min_cluster_size=3))])
```

## 1.7 Mini Batch KMeans

```
[124]: mb_kmeans = MiniBatchKMeans(n_clusters=3, max_iter=1)
```

```
[125]: clf = fit_model(algorithm=mb_kmeans, data=(X_small, y_small))
```

Started Training MiniBatchKMeans on X: (10937,) y: (10937,)

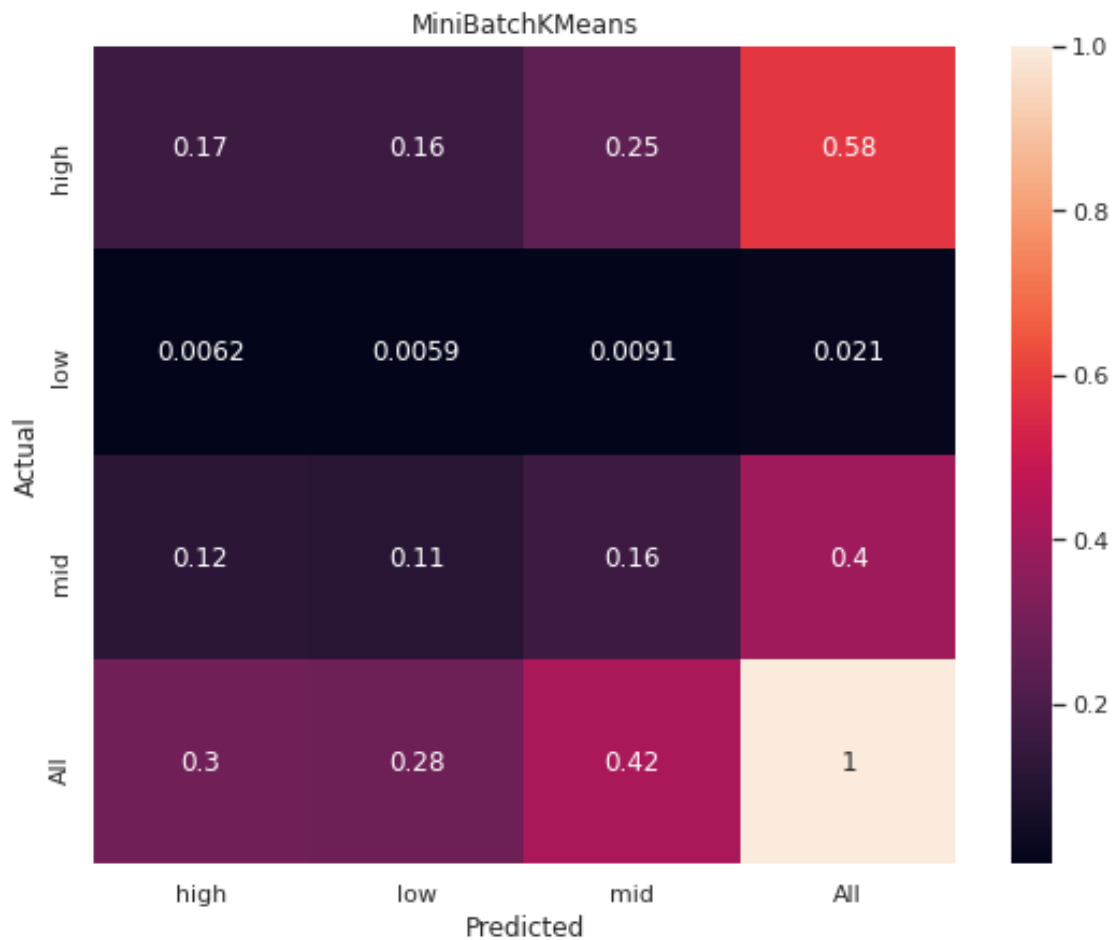
Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for MiniBatchKMeans

	precision	recall	f1-score	support
low	0.02	0.28	0.04	231
mid	0.39	0.41	0.40	4333
high	0.57	0.29	0.38	6373
accuracy			0.34	10937
macro avg	0.33	0.33	0.27	10937
weighted avg	0.49	0.34	0.38	10937

Accuracy Score: 0.3378440157264332

Balanced Accuracy Score: 0.32632201462411303



Trained MiniBatchKMeans in 217.56095910072327s

```
[126]: clf
```

```
[126]: Pipeline(steps=[('mean_embeddings',  
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at  
                        0x7ff253d3b6d8>)),  
                      ('reduce_dim', TruncatedSVD(n_components=50)),  
                      ('clusterer', MiniBatchKMeans(max_iter=1, n_clusters=3))])
```

```
[ ]:
```