

# Assignment

Course Code	CSC402A
Course Name	Data Mining
Programme	B.Tech
Department	CSE
Faculty	FET

Name of the Student	Satyajit Ghana
Reg. No.	17ETCS002159
Semester/Year	07/2021
Course Leader(s)	Prof. Mohan Kumar

# Declaration Sheet

Student Name	Satyajit Ghana		
Reg. No	17ETCS002159		
Programme	B.Tech	Semester/Year	07/2021
Course Code	CSC402A		
Course Title	Data Mining		
Course Date		to	
Course Leader	Prof. Mohan Kumar		
<p><b>Declaration</b></p> <p>The assignment submitted herewith is a result of my own investigations and that I have conformed to the guidelines against plagiarism as laid out in the Student Handbook. All sections of the text and results, which have been obtained from other sources, are fully referenced. I understand that cheating and plagiarism constitute a breach of University regulations and will be dealt with accordingly.</p>			
Signature of the Student		Date	
Submission date stamp (by Examination & Assessment Section)			
Signature of the Course Leader and date		Signature of the Reviewer and date	

# Contents

Declaration Sheet	ii
Contents	iii
1 Question 1	4
1.1 Introduction	4
1.2 Supervised Learning	7
1.3 Unsupervised Learning	9
1.4 Comparative Analysis and Conclusion	11
1.4.1 Conclusion	13
2 Appendix	14

# 1 Question 1

Solution to Question No. 1 Part B

## 1.1 Introduction

A few more pre-processing steps were performed on the dataset, to clean it, which will then be used by the learning algorithms.

Columns like [user\_id, isbn, city, state] are dropped, since they don't contribute to the rating classification, and they don't really tell a lot about rating of the book (not much of information gain), also something to note (since we are considering the country) column, that most of the users (~70%) are from usa, which may/may not contribute to lot to the accuracy of the classification.

```
[ ] 1 dataset.head()
```

	age	book_rating	book_title	book_author	year_of_publication		publisher	country
0	34	5	Clara Callan	Richard Bruce Wright	2001	HarperFlamingo	Canada	canada
2	30	8	Clara Callan	Richard Bruce Wright	2001	HarperFlamingo	Canada	canada
4	34	9	Clara Callan	Richard Bruce Wright	2001	HarperFlamingo	Canada	canada
5	34	8	Clara Callan	Richard Bruce Wright	2001	HarperFlamingo	Canada	canada
6	34	9	Clara Callan	Richard Bruce Wright	2001	HarperFlamingo	Canada	canada

The target column is **book\_rating**, now we have 10 classes for **book\_rating**, but to make the task a little simpler we will be labelling the **book\_rating** into 3 classes, namely [low, mid, high].

The stats for them are,

high	mid	low
58.482%	39.248%	2.269%

```
[ ] 1 bins = [0, 3, 7, 10]
    2 names = ['low', 'mid', 'high']
    3
    4 dataset['book_rating'] = pd.cut(dataset['book_rating'], bins, labels=names)
```

```
[ ] 1 dataset.head()
```

	age	book_rating	book_title	book_author	year_of_publication	publisher	country
0	34	mid	Clara Callan	Richard Bruce Wright	2001	HarperFlamingo Canada	canada
2	30	high	Clara Callan	Richard Bruce Wright	2001	HarperFlamingo Canada	canada
4	34	high	Clara Callan	Richard Bruce Wright	2001	HarperFlamingo Canada	canada
5	34	high	Clara Callan	Richard Bruce Wright	2001	HarperFlamingo Canada	canada
6	34	high	Clara Callan	Richard Bruce Wright	2001	HarperFlamingo Canada	canada

```
[ ] 1 dataset['book_rating'].value_counts()
```

```
high    213208
mid     143087
low       8275
Name: book_rating, dtype: int64
```

To perform any kind of learning on this dataset we would have to convert the categorical values into numerical, to do this there are two options

- OneHotEncoder
- LabelEncoder

Label Encoder is used for those categorical values which have a “ordering” associated with them, for example the temperature can be cold, warm or hot, we can assign the values 0, 1, 2 to them, that makes sense right? Yes! But now let’s say the categorical values are Satyajit, Mohan, Ram how can a number be assigned to them? We cannot assign 0, 1, 2 to them, because we don’t how to order them, to deal with this problem we use OneHotEncoder, which simply assigns them [1, 0, 0], [0, 1, 0] and [0, 0, 1] which are basically indicating presence, of [Satyajit, Mohan, Ram] and there is no ordering involved here.

But looking at our dataset, we have a lot of data, and the since we are considering the book name, author, publisher all as categorical value, they will have quite a lot of categories, which

means that when we do one hot encoding, there will be a lot of columns generating, which will be very sparse ofcourse, but this means that we will be restricted by memory consumption.

```
[12] 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

```
[13] 1 numeric_features = ['age', 'year_of_publication']  
2 numeric_transformer = Pipeline(steps=[  
3     ('imputer', SimpleImputer(strategy='median')),  
4     ('scaler', StandardScaler())])
```

```
[14] 1 categorical_features = ['book_title', 'book_author', 'publisher', 'country']  
2 categorical_transformer = OneHotEncoder(handle_unknown='ignore')
```

```
[15] 1 preprocessor = ColumnTransformer(  
2     transformers=[  
3         ('num', numeric_transformer, numeric_features),  
4         ('cat', categorical_transformer, categorical_features)])
```

```
[ ] 1 dummy = preprocessor.fit_transform(X_train)  
2 dummy.shape
```

```
(291656, 177473)
```

If we were to store all the all data, and each cell occupying 4 bytes, this is how much space we would require

```
[ ] 1 f'{{(dummy.shape[0] * dummy.shape[1] * 4) // 1024 // 1024 // 1024}} giga bytes'  
  
'192 giga bytes'
```

From the above illustration, we can see if we try to save this data, we would approximately require 192 GB, and that is just the training set. (I did try it on a 128GB RAM machine, and I got out-of-memory error). We'll later discuss some of the solutions to this.

## 1.2 Supervised Learning

Supervised Learning was performed on the cleaned dataset, with OneHotEncoding and StandardScaler with the following algorithms

- Decision Tree
- SVC (Support Vector Classification)
- Passive Aggressive
- Ridge Classifier
- AdaBoost
- LinearSVC
- MLP (Multi Layer Perceptron)
- RandomForest

Refer to the notebook attached at the end for more information of accuracy, confusion matrix, f1 score, etc., here we will compare the accuracies and time taken to train for all the models.

Algorithm	F1 Score	Balanced Accuracy	Avg Precision	Avg Recall	Time Taken (s)
DecisionTree	0.60	0.36	0.58	0.60	146
SVC	NULL	NULL	NULL	NULL	>1000
PassiveAggressive	0.55	0.37	0.55	0.55	8
Ridge	0.58	0.37	0.56	0.59	22
AdaBoost	0.59	0.34	0.57	0.59	50
LinearSVC	0.58	0.37	0.56	0.59	200
<b>MLP</b>	<b>0.61</b>	<b>0.38</b>	<b>0.60</b>	<b>0.61</b>	630
RandomForest	0.58	0.33	0.60	0.59	537

All of the above models were trained by using Imputation, and Standard Scaler, but we can see and vary the scaler to see how it affects the accuracy, to do that we consider the following scalers, and since Passive Aggressive algorithm was the faster, we will stick with that,

- StandardScaler
- MinMaxScaler
- RobustScaler
- MaxAbsScaler
- PowerTransformer
- QuantileTransformer
- Normalizer
- UnScaled (Original Data)

Scaler	F1 Score	Balanced Accuracy	Avg Precision	Avg Recall
StandardScaler	0.561	<b>0.378</b>	0.55	0.56
MinMaxScaler	0.541	0.369	0.55	0.54
<b>MaxAbsScaler</b>	<b>0.588</b>	0.359	<b>0.55</b>	<b>0.59</b>
RobustScaler	0.576	0.366	0.54	0.58
PowerTransform (Yeo-Johnson)	0.525	0.374	0.55	0.53
QuantileTransform	0.519	0.371	0.54	0.52
Normalizer (L2)	0.570	0.366	0.54	0.57
UnScaled	0.583	0.364	0.55	0.58



### 1.3 Unsupervised Learning

In the case of unsupervised algorithms, most of them crashed on the original data, since the data was very sparse, and a lot of memory was required.

Learning from the problems faced in supervised learning, one solution to deal with text data is to perform NLP (Natural Language Processing) on the text itself, i.e. we can use pretrained word embeddings to encode the book names, authors, publishers, country into a dense matrix, which can be easily done by using spacy language models.

All of the columns in the dataset was concatenated to form a single sentence, we will later encode this sentence and try to apply clustering algorithm to predict rating.

```
[8] 1 c_dataset = dataset["book_title"].astype(str) + " " + \
2      dataset["book_author"].astype(str) + " " + \
3      dataset["publisher"].astype(str) + " " + \
4      dataset["year_of_publication"].astype(str) + " " + \
5      dataset["age"].astype(str) + " " + \
6      dataset["country"].astype(str)
7
```

```
[9] 1 for ex in c_dataset.sample(frac=0.2)[:5]:
2     print(ex)
```

```
The Vanishing Vampire (The Accidental Monsters , No 1) David Lubar Scholastic 1997.0 12.0 usa
REMEMBER ME Mary Higgins Clark Simon & Schuster 1994.0 34.0 usa
The Mummy or Ramses the Damned Anne Rice Ballantine Books 1991.0 22.0 usa
Bittersweet Rain Sandra Brown Warner Books 2000.0 34.0 usa
Arthur Stephen R. Lawhead Zondervan Publishing Company 1996.0 21.0 usa
```

Also learning from the fact that I don't have enough time to fit a model on this large (300,000) data, we can simplify evaluation of algorithms by sampling a fraction of the original dataset. Even on a really good machine (i7-5930K, 128GB RAM), the model takes a lot of time to fit.

The following algorithms were used,

- KMedoids Clustering

- Agglomerative Clustering
- DBSCAN
- KMeans
- HDBSCAN
- MiniBatchKMeans

Algorithm	F1 Score	Balanced Accuracy	Avg Precision	Avg Recall	Time Taken (s)
KMedoids	0.30	0.31	0.50	0.31	223
Agglomerative	0.27	0.30	0.48	0.27	119
DBSCAN	<b>0.58</b>	<b>0.33</b>	0.42	<b>0.58</b>	<b>116</b>
KMeans	0.31	0.33	0.50	0.32	222
<b>HDBSCAN</b>	<b>0.58</b>	<b>0.33</b>	<b>0.54</b>	<b>0.58</b>	121
MiniBatch KMeans	0.33	0.32	0.49	0.34	217

To see the classification report, and the confusion matrix, see appendix.

## 1.4 Comparative Analysis and Conclusion

If we were to compare the supervised and unsupervised models, they have pretty much comparable accuracies (~58-61%), but one thing that cannot be compared is the execution time, since we used clustering algorithms, they require heavy memory usage, and a lot of processing time.

The highest accuracy in Supervised Learning was achieved by MLP (Multi-Layer Perceptron) of 61%, and in Unsupervised Learning, highest was achieved by HDBSCAN of 58%. And in terms of Scaler, it was observed that MaxAbsScaler gave the highest accuracy compared to other scalers.

Something even more important to address is why the accuracy is so low, the **"high"** class of `book_rating` has 58% total records, and this means that if we were to classify any input as **"high"** we have a probability of 0.58 to be correct, this is exactly like a **OneR** learner, but the reason why we are observing such bad accuracy is because there isn't much correlation between the attributes and the target value, this dataset was never meant to do something like this, it was more towards recommending books to a new user with a given past book purchase history, predicting rating based on country, book title, publisher just does not make sense, just based on the book title, how could a book be rated better? One thing that could make sense is the author, there are certain authors that have a good reputation and they have a higher chance of getting a good rating. But in this assignment, we have considered all the attributes, which could have been a mistake, one lesson learnt is that, grabbing all the available data can have negative effects to the accuracy.

But there is another issue, if we see the number of categorical values for each of the categorical columns,

```
[136] 1 dataset['book_title'].cat.categories.shape
      (132033,)

[139] 1 dataset['book_author'].cat.categories.shape
      (60652,)

[140] 1 dataset['publisher'].cat.categories.shape
      (11311,)

[141] 1 dataset['country'].cat.categories.shape
      (51,)
```

Something we can clearly observe is that, that is a lot of categories, for specially the book\_title, this means that if we use OneHotEncoding, we will get those many columns added to the dataset. This does make it really difficult to train a model, scipy has a csr representation (Compressed Sparse Row Matrix) which does help to a lot of extend, but still this is the biggest issue that has to be addressed.

We can take inspiration from text classification models and NLP (Natural Language Processing), what is done for it, is that each of the words in dictionary is a sparse matrix, and then a model is used to convert this sparse matrix into a dense matrix, for example lets say the word "the" is represented as [0, 1, 0, 0, 0, 0, 0, 0, 0, 0], if we use a word embedding model, such as from spacy, we would get something like [0.45673, 0.67395, 0.98729] which is a much better representation as before, this matrix can then be used to train. This inspiration was tried in the clustering algorithms in 1.3, and it quite did not produce good results. But I would blame the dataset for that. On top of giving the dense matrix, we can further reduce the dimensionality by using SVD (Singular Value Decomposition).

I had a hunch that if I were to use GPT-2 or BERT, the model would have been really great (although it might overfit, but still), but due to time constrains, it wasn't possible, but it would be a good exploration domain.

### 1.4.1 Conclusion

Never rely on SciKit Learn, the algorithms are really slow (except K-Mean) and are not quite optimized enough, instead C++ backend algorithms give much better speed.

Also, when a classification model has to be made, in which there is a lot of attributes in text in them, it would be better, to just convert those specific attributes to a dense matrix using a pretrained word embedding model, or an even better improvement can be to use **n-gram** models.

To conclude, we were not really successful in creating a good classification model for the given dataset, but we were definitely able to compare different supervised models, unsupervised models and numeric scaling techniques, overall, the outcome of this course was covered.

## 2 Appendix

---

# BookReview-Classification-Clean

January 20, 2021

```
[ ]: import gdown

url = 'https://drive.google.com/uc?id=1PL13wgXLfXcsrKKNuVIAJdGXrIqv2mv'
output = 'book_crossing.cleaned.csv'
gdown.download(url, output, quiet=False)
```

Downloading...

From: <https://drive.google.com/uc?id=1PL13wgXLfXcsrKKNuVIAJdGXrIqv2mv>

To: /content/book\_crossing.cleaned.csv

44.9MB [00:01, 40.7MB/s]

```
[ ]: 'book_crossing.cleaned.csv'
```

```
[ ]: %matplotlib inline

import scipy
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

sns.set()
palette = sns.color_palette("icefire")

plt.style.use('ggplot')

sns.set_context("talk")
```

## 1 BookCrossing - Cleaning

```
[ ]: dataset = pd.read_csv('book_crossing.cleaned.csv')
```

```
[ ]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 383849 entries, 0 to 383848
```

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	user_id	383849 non-null	int64
1	age	383849 non-null	int64
2	isbn	383849 non-null	object
3	book_rating	383849 non-null	int64
4	book_title	383849 non-null	object
5	book_author	383849 non-null	object
6	year_of_publication	383849 non-null	int64
7	publisher	383849 non-null	object
8	city	375223 non-null	object
9	state	371248 non-null	object
10	country	366406 non-null	object

dtypes: int64(4), object(7)

memory usage: 32.2+ MB

We won't be considering city, state, because they don't really tell a lot of the rating of a book, but also most of the users (~70%) are from usa (which may not contribute a lot to accuracy of classification, but we'll consider it), and the location is related to the user, and not the book directly, we'll also be dropping isbn, user\_id, since they don't contribute to classification of rating

```
[ ]: dataset = dataset.drop(['user_id', 'isbn', 'city', 'state'], axis=1)
```

```
[ ]: f'Dataset Shape : {dataset.shape}'
```

```
[ ]: 'Dataset Shape : (383849, 7)'
```

```
[ ]: dataset.dropna(inplace=True)
```

```
[ ]: f'Dataset Shape after dropping NA: {dataset.shape}'
```

```
[ ]: 'Dataset Shape after dropping NA: (366406, 7)'
```

```
[ ]: dataset.head()
```

```
[ ]: 
```

	age	book_rating	...		publisher	country
0	34	5	...	HarperFlamingo	Canada	canada
2	30	8	...	HarperFlamingo	Canada	canada
4	34	9	...	HarperFlamingo	Canada	canada
5	34	8	...	HarperFlamingo	Canada	canada
6	34	9	...	HarperFlamingo	Canada	canada

[5 rows x 7 columns]

```
[ ]: dataset.describe().T
```

```
[ ]: 
```

	count	mean	std	...	50%	75%
max						
age	366406.0	35.860998	10.448608	...	34.0	40.0
100.0						
book_rating	366406.0	7.635975	1.836354	...	8.0	9.0
10.0						



```
year_of_publication 366406.0 1995.670314 7.397156 ... 1997.0 2001.0
2006.0
```

```
[3 rows x 8 columns]
```

We'll remove the rows which have a country which has value count <= 50

```
[ ]: dataset = dataset.groupby('country').filter(lambda x: len(x) > 50)
```

```
[ ]: dataset.describe().T
```

```
[ ]:
      count      mean      std  ...    50%    75%
max
age      364570.0    35.867227  10.447887  ...   34.0   40.0
100.0
book_rating      364570.0     7.636709   1.835857  ...    8.0    9.0
10.0
year_of_publication 364570.0 1995.667164   7.400552  ...  1997.0  2001.0
2006.0
```

```
[3 rows x 8 columns]
```

```
[ ]: f'Dataset Shape : {dataset.shape}'
```

```
[ ]: 'Dataset Shape : (364570, 7)'
```

```
[ ]: f'Column Names: {dataset.columns.to_list()}'
```

```
[ ]: "Column Names: ['age', 'book_rating', 'book_title', 'book_author',
'year_of_publication', 'publisher', 'country']"
```

```
[ ]: dataset['book_rating'].value_counts()
```

```
[ ]: 8      87090
      10     68038
      7     63036
      9     58080
      5     42988
      6     29943
      4      7120
      3     4746
      2     2198
      1     1331
```

```
Name: book_rating, dtype: int64
```

We'll now convert the rating into classification categories

```
[ ]: bins = [0, 3, 7, 10]
```

```
names = ['low', 'mid', 'high']
```

```
dataset['book_rating'] = pd.cut(dataset['book_rating'], bins, labels=names)
```

```
[ ]: dataset.head()
```

```
[ ]:   age book_rating ...           publisher country
      0    34         mid ...  HarperFlamingo Canada  canada
      2    30         high ...  HarperFlamingo Canada  canada
      4    34         high ...  HarperFlamingo Canada  canada
      5    34         high ...  HarperFlamingo Canada  canada
      6    34         high ...  HarperFlamingo Canada  canada
```

[5 rows x 7 columns]

```
[ ]: dataset['book_rating'].value_counts()
```

```
[ ]: high    213208
      mid     143087
      low      8275
      Name: book_rating, dtype: int64
```

```
[ ]: dataset.to_csv('book_crossing.classification.cleaned.csv', index=False)
```

```
[ ]:
```

# BookReview-Classification-Supervised

January 20, 2021

```
[1]: ! pip install --upgrade scikit-learn
```

## 1 Book Crossing - Classification

```
[1]: %matplotlib inline

import scipy
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

sns.set()
palette = sns.color_palette("icefire")

plt.style.use('ggplot')

sns.set_context("talk")

[2]: dataset = pd.read_csv('book_crossing.classification.cleaned.csv')

[3]: dataset['age'] = dataset['age'].astype(np.float64)
dataset['book_rating'] = dataset['book_rating'].astype('category')
dataset['book_title'] = dataset['book_title'].astype('category')
dataset['book_author'] = dataset['book_author'].astype('category')
dataset['year_of_publication'] = dataset['year_of_publication'].astype(np.
    ↳float64)
dataset['publisher'] = dataset['publisher'].astype('category')
dataset['country'] = dataset['country'].astype('category')

[4]: dataset.head()

[4]:   age book_rating  ...      publisher country
0  34.0         mid  ...  HarperFlamingo  Canada  canada
1  30.0         high  ...  HarperFlamingo  Canada  canada
2  34.0         high  ...  HarperFlamingo  Canada  canada
3  34.0         high  ...  HarperFlamingo  Canada  canada
```

```
4  34.0          high ... HarperFlamingo Canada  canada
```

```
[5 rows x 7 columns]
```

```
[5]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 364570 entries, 0 to 364569
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	age	364570 non-null	float64
1	book_rating	364570 non-null	category
2	book_title	364570 non-null	category
3	book_author	364570 non-null	category
4	year_of_publication	364570 non-null	float64
5	publisher	364570 non-null	category
6	country	364570 non-null	category

```
dtypes: category(5), float64(2)
```

```
memory usage: 19.1 MB
```

```
[6]: from sklearn import set_config
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder, LabelEncoder, \
    \->MinMaxScaler
from sklearn.svm import SVC, LinearSVC
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, plot_confusion_matrix, \
    \->confusion_matrix, accuracy_score, balanced_accuracy_score
from sklearn.neural_network import MLPClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
from sklearn.linear_model import RidgeClassifier, PassiveAggressiveClassifier

from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import minmax_scale
from sklearn.preprocessing import MaxAbsScaler
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import RobustScaler
from sklearn.preprocessing import Normalizer
from sklearn.preprocessing import QuantileTransformer
from sklearn.preprocessing import PowerTransformer

set_config(display='diagram')
```

```
[7]: X, y = dataset.drop('book_rating', axis=1), dataset['book_rating']
```

```
[8]: target_names = ['low', 'mid', 'high']
```

```
[9]: X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 364570 entries, 0 to 364569
Data columns (total 6 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   364570 non-null float64
 1   book_title            364570 non-null category
 2   book_author           364570 non-null category
 3   year_of_publication   364570 non-null float64
 4   publisher              364570 non-null category
 5   country                364570 non-null category
dtypes: category(4), float64(2)
memory usage: 18.8 MB
```

```
[10]: X.head()
```

```
[10]:    age  book_title  ...  publisher  country
0  34.0  Clara Callan  ...  HarperFlamingo  Canada  canada
1  30.0  Clara Callan  ...  HarperFlamingo  Canada  canada
2  34.0  Clara Callan  ...  HarperFlamingo  Canada  canada
3  34.0  Clara Callan  ...  HarperFlamingo  Canada  canada
4  34.0  Clara Callan  ...  HarperFlamingo  Canada  canada
```

```
[5 rows x 6 columns]
```

```
[11]: y.head()
```

```
[11]: 0    mid
1    high
2    high
3    high
4    high
Name: book_rating, dtype: category
Categories (3, object): ['high', 'low', 'mid']
```

```
[12]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=0)
```

```
[13]: numeric_features = ['age', 'year_of_publication']
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])
```

```
[14]: categorical_features = ['book_title', 'book_author', 'publisher', 'country']
categorical_transformer = OneHotEncoder(handle_unknown='ignore')
```

```
[15]: preprocessor = ColumnTransformer(  
      transformers=[  
          ('num', numeric_transformer, numeric_features),  
          ('cat', categorical_transformer, categorical_features)])
```

```
[24]: dummy = preprocessor.fit_transform(X_train)  
      dummy.shape
```

```
[24]: (291656, 177473)
```

If we were to store all the all data, and each cell occupying 4 bytes, this is how much space we would require

```
[30]: f'{{(dummy.shape[0] * dummy.shape[1] * 4) // 1024 // 1024 // 1024}} giga bytes'
```

```
[30]: '192 giga bytes'
```

## 1.1 Supervised Models

- DecisionTreeClassifier
- SVC
- LinearSVC
- KNeighborsClassifier
- PassiveAggressiveClassifier
- AdaBoostClassifier
- GaussianProcessClassifier

```
[16]: from time import time  
  
def fit_model(algorithm, data, preprocessor):  
  
    t1 = time()  
  
    X_train, X_test, y_train, y_test = data  
  
    clf = Pipeline(steps=[('preprocessor', preprocessor),  
                          ('classifier', algorithm)])  
  
    print(f'\nStarted Training {algorithm.__class__.__name__} on X_train:   
→{X_train.shape} y_train: {y_train.shape}')  
  
    # train the model  
    clf.fit(X_train, y_train)  
  
    print(f'\nEvaluating model on X_test: {X_test.shape} y_test: {y_test.  
→shape}')  
  
    # test the model  
    y_true = y_test.copy()  
    y_pred = clf.predict(X_test)
```

```

# get the classification report
print(f"\nClassification Report for {algorithm.__class__.__name__}")
print(classification_report(y_true, y_pred, target_names=target_names,
→labels=target_names))

acc_score = accuracy_score(y_true, y_pred)
bal_score = balanced_accuracy_score(y_true, y_pred)

print(f"\nAccuracy Score: {acc_score}")
print(f"Balanced Accuracy Score: {bal_score}")

print()
# show the confusion matrix
cmmat_table = pd.DataFrame({'y_true': y_true, 'y_pred': y_pred})
conmat = pd.crosstab(cmmat_table.y_true, cmmat_table.y_pred,
→rownames=['Actual'], colnames=['Predicted'], margins=True, normalize='all')
ax = plt.axes()
sns.set(rc={'figure.figsize':(9, 7)})
sns.heatmap(conmat, annot=True, ax=ax)
ax.set_title(f'{algorithm.__class__.__name__}')
plt.show()
print()

t2 = time()

print(f'Trained {algorithm.__class__.__name__} in {(t2 - t1)}s')

return clf

```

## 1.2 DecisionTree Classifier

```
[33]: dtc = DecisionTreeClassifier(max_depth=100)
```

```
[34]: clf = fit_model(algorithm=dtc, data=(X_train, X_test, y_train, y_test),
→preprocessor=preprocessor)
```

Started Training DecisionTreeClassifier on X\_train: (291656, 6) y\_train: (291656,)

Evaluating model on X\_test: (72914, 6) y\_test: (72914,)

```


Classification Report for DecisionTreeClassifier
precision    recall  f1-score   support

```

low	0.19	0.03	0.05	1603
mid	0.55	0.17	0.25	28483
high	0.61	0.91	0.73	42828
accuracy			0.60	72914
macro avg	0.45	0.37	0.35	72914
weighted avg	0.58	0.60	0.53	72914

Accuracy Score: 0.6010368379186439

Balanced Accuracy Score: 0.36884514138535285



Trained DecisionTreeClassifier in 146.20573687553406s

[35]: clf



```
[35]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer(strategy='median')),
                                                                              StandardScaler())]),
                                                         ('cat',
                                                                              OneHotEncoder(handle_unknown='ignore'),
                                                                              ['book_title', 'book_author',
                                                                              'publisher', 'country'])])),
                        ('classifier', DecisionTreeClassifier(max_depth=100))])
```

### 1.3 SVC (Support Vector Classification)

```
[31]: svc = SVC(gamma='auto', C=1.0, kernel='rbf')
[:]: clf = fit_model(algorithm=svc, data=(X_train, X_test, y_train, y_test),
                    ↪preprocessor=preprocessor)
```

Started Training SVC on X\_train: (291656, 6) y\_train: (291656,)

```
[]: clf
```

### 1.4 Passive Aggressive Classifier

```
[17]: pac = PassiveAggressiveClassifier(early_stopping=True, verbose=1, n_jobs=-1)
[20]: clf = fit_model(algorithm=pac, data=(X_train, X_test, y_train, y_test),
                    ↪preprocessor=preprocessor)
```

Started Training PassiveAggressiveClassifier on X\_train: (291656, 6) y\_train: (291656,)

[Parallel(n\_jobs=-1)]: Using backend ThreadingBackend with 2 concurrent workers.

-- Epoch 1-- Epoch 1

Norm: 46.21, NNZs: 53157, Bias: -0.504273, T: 262490, Avg. loss: 0.105049  
Total training time: 0.11 seconds.

-- Epoch 2

Norm: 119.01, NNZs: 149844, Bias: 0.084396, T: 262490, Avg. loss: 0.905595  
Total training time: 0.18 seconds.

-- Epoch 2

Norm: 66.08, NNZs: 70363, Bias: -0.521754, T: 524980, Avg. loss: 0.086252  
 Total training time: 0.29 seconds.  
 -- Epoch 3  
 Norm: 164.90, NNZs: 161910, Bias: 0.072078, T: 524980, Avg. loss: 0.789480  
 Total training time: 0.36 seconds.  
 -- Epoch 3  
 Norm: 79.79, NNZs: 79611, Bias: -0.583871, T: 787470, Avg. loss: 0.077301  
 Total training time: 0.46 seconds.  
 -- Epoch 4  
 Norm: 198.89, NNZs: 164254, Bias: 0.097328, T: 787470, Avg. loss: 0.731145  
 Total training time: 0.54 seconds.  
 -- Epoch 4  
 Norm: 90.21, NNZs: 85297, Bias: -0.646251, T: 1049960, Avg. loss: 0.072339  
 Total training time: 0.63 seconds.  
 -- Epoch 5  
 Norm: 225.85, NNZs: 164894, Bias: 0.119736, T: 1049960, Avg. loss: 0.696331  
 Total training time: 0.72 seconds.  
 -- Epoch 5  
 Norm: 98.61, NNZs: 89244, Bias: -0.634484, T: 1312450, Avg. loss: 0.070211  
 Total training time: 0.80 seconds.  
 -- Epoch 6  
 Norm: 248.75, NNZs: 165157, Bias: 0.069723, T: 1312450, Avg. loss: 0.674325  
 Total training time: 0.89 seconds.  
 -- Epoch 6  
 Norm: 105.81, NNZs: 92113, Bias: -0.705083, T: 1574940, Avg. loss: 0.067957  
 Total training time: 0.99 seconds.  
 Convergence after 6 epochs took 1.04 seconds  
 -- Epoch 1  
 Norm: 268.11, NNZs: 165276, Bias: 0.086452, T: 1574940, Avg. loss: 0.659386  
 Total training time: 1.10 seconds.  
 Convergence after 6 epochs took 1.16 seconds  
 Norm: 117.95, NNZs: 148735, Bias: -0.110863, T: 262490, Avg. loss: 0.897951  
 Total training time: 0.15 seconds.  
 -- Epoch 2  
 Norm: 163.59, NNZs: 161578, Bias: -0.108644, T: 524980, Avg. loss: 0.784290  
 Total training time: 0.30 seconds.  
 -- Epoch 3  
 Norm: 197.51, NNZs: 164105, Bias: -0.107356, T: 787470, Avg. loss: 0.725686  
 Total training time: 0.45 seconds.  
 -- Epoch 4  
 Norm: 224.64, NNZs: 164845, Bias: -0.101451, T: 1049960, Avg. loss: 0.692328  
 Total training time: 0.59 seconds.  
 -- Epoch 5  
 Norm: 247.11, NNZs: 165150, Bias: -0.105241, T: 1312450, Avg. loss: 0.670131  
 Total training time: 0.73 seconds.  
 -- Epoch 6  
 Norm: 266.48, NNZs: 165274, Bias: -0.121605, T: 1574940, Avg. loss: 0.654409  
 Total training time: 0.89 seconds.

```
-- Epoch 7
Norm: 283.55, NNZs: 165349, Bias: -0.141001, T: 1837430, Avg. loss: 0.642702
Total training time: 1.03 seconds.
-- Epoch 8
Norm: 298.48, NNZs: 165405, Bias: -0.106583, T: 2099920, Avg. loss: 0.634828
Total training time: 1.18 seconds.
-- Epoch 9
Norm: 312.23, NNZs: 165423, Bias: -0.145596, T: 2362410, Avg. loss: 0.626431
Total training time: 1.33 seconds.
Convergence after 9 epochs took 1.38 seconds
```

```
[Parallel(n_jobs=-1)]: Done    3 out of    3 | elapsed:    2.5s finished
```


```
Evaluating model on X_test: (72914, 6) y_test: (72914,)
```

```
Classification Report for PassiveAggressiveClassifier
           precision    recall  f1-score   support

     low           0.04       0.09       0.05         1603
     mid           0.47       0.29       0.36        28483
     high           0.62       0.75       0.68        42828

 accuracy                   0.55         72914
 macro avg           0.38       0.37       0.36         72914
weighted avg           0.55       0.55       0.54         72914
```

```
Accuracy Score: 0.5531722302987081
Balanced Accuracy Score: 0.37477427261674867
```



Trained PassiveAggressiveClassifier in 8.378920316696167s

[21]: clf

```
[21]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer(strategy='median')),
                                                                              ('scaler',
                                                                              StandardScaler()))],
                                                         ['age',
                                                                              'year_of_publication']),
                                                         ('cat',
                                                         OneHotEncoder(handle_unknown='ignore'),
                                                         ['book_title', 'book_author',
                                                                              'publisher', 'country'])])),
                      ('classifier',
                       PassiveAggressiveClassifier(early_stopping=True, n_jobs=-1,
                                                     verbose=1))])
```

## 1.5 Ridge Classifier

```
[26]: rc = RidgeClassifier(tol=1e-2, solver="auto")
```

```
[27]: clf = fit_model(algorithm=rc, data=(X_train, X_test, y_train, y_test),  
    ↪preprocessor=preprocessor)
```

Started Training RidgeClassifier on X\_train: (291656, 6) y\_train: (291656,)


Evaluating model on X\_test: (72914, 6) y\_test: (72914,)

Classification Report for RidgeClassifier

	precision	recall	f1-score	support
low	0.06	0.01	0.02	1603
mid	0.49	0.35	0.41	28483
high	0.63	0.77	0.69	42828
accuracy			0.59	72914
macro avg	0.39	0.38	0.37	72914
weighted avg	0.56	0.59	0.57	72914

Accuracy Score: 0.5888855363853307

Balanced Accuracy Score: 0.3765274882659582



Trained RidgeClassifier in 22.259966611862183s

[28]: clf

```
[28]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer(strategy='median')),
                                                                              ('scaler',
                                                                              StandardScaler()))],
                                                         ['age',
                                                                              'year_of_publication'])),
                        ('cat',
                        OneHotEncoder(handle_unknown='ignore'),
                        ['book_title', 'book_author',
                        'publisher', 'country'])))],
          ('classifier', RidgeClassifier(tol=0.01))])
```

## 1.6 AdaBoost Classifier

```
[29]: abc = AdaBoostClassifier()
```

```
[30]: clf = fit_model(algorithm=abc, data=(X_train, X_test, y_train, y_test),  
    ↳preprocessor=preprocessor)
```

Started Training AdaBoostClassifier on X\_train: (291656, 6) y\_train: (291656,)


Evaluating model on X\_test: (72914, 6) y\_test: (72914,)

Classification Report for AdaBoostClassifier

	precision	recall	f1-score	support
low	0.40	0.01	0.02	1603
mid	0.55	0.05	0.09	28483
high	0.59	0.97	0.74	42828
accuracy			0.59	72914
macro avg	0.51	0.34	0.28	72914
weighted avg	0.57	0.59	0.47	72914

Accuracy Score: 0.5919165043750172

Balanced Accuracy Score: 0.34478823998711056



Trained AdaBoostClassifier in 50.97685408592224s

[31]: clf

```
[31]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer(strategy='median')),
                                                                              ('scaler',
                                                                              StandardScaler()))],
                                                         ['age',
                                                                              'year_of_publication']),
                                                         ('cat',
                                                         OneHotEncoder(handle_unknown='ignore'),
                                                         ['book_title', 'book_author',
                                                                              'publisher', 'country'])])),
                        ('classifier', AdaBoostClassifier())])
```



## 1.7 Linear SVC

```
[28]: lsvc = LinearSVC(verbose=1)
[29]: clf = fit_model(algorithm=lsvc, data=(X_train, X_test, y_train, y_test),
    ↳preprocessor=preprocessor)
```

Started Training LinearSVC on X\_train: (291656, 6) y\_train: (291656,)

/usr/local/lib/python3.6/dist-packages/sklearn/svm/\_base.py:986:  
ConvergenceWarning: Liblinear failed to converge, increase the number of  
iterations.

"the number of iterations.", ConvergenceWarning)


Evaluating model on X\_test: (72914, 6) y\_test: (72914,)

Classification Report for LinearSVC

	precision	recall	f1-score	support
low	0.06	0.01	0.02	1603
mid	0.49	0.35	0.41	28483
high	0.63	0.76	0.69	42828
accuracy			0.59	72914
macro avg	0.39	0.38	0.37	72914
weighted avg	0.56	0.59	0.57	72914

Accuracy Score: 0.5862660120141536

Balanced Accuracy Score: 0.3763315404262985



Trained LinearSVC in 199.0021414756775s

[30]: clf

```
[30]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer(strategy='median')),
                                                                              ('scaler',
                                                                              StandardScaler()))],
                                                         ['age',
                                                                              'year_of_publication'])),
                        ('cat',
                        OneHotEncoder(handle_unknown='ignore'),
                        ['book_title', 'book_author',
                        'publisher', 'country'])))],
        ('classifier', LinearSVC(verbose=1))])
```

## 1.8 MLP Classifier

```
[26]: mlp = MLPClassifier(alpha=0.001, max_iter=1, verbose=1)
```

```
[27]: clf = fit_model(algorithm=mlp, data=(X_train, X_test, y_train, y_test),  
    ↳preprocessor=preprocessor)
```

Started Training MLPClassifier on X\_train: (291656, 6) y\_train: (291656,)  
Iteration 1, loss = 0.75835927

Evaluating model on X\_test: (72914, 6) y\_test: (72914,)


```
/usr/local/lib/python3.6/dist-  
packages/sklearn/neural_network/_multilayer_perceptron.py:617:  
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1) reached and the  
optimization hasn't converged yet.  
% self.max_iter, ConvergenceWarning)
```

Classification Report for MLPClassifier

	precision	recall	f1-score	support
low	1.00	0.00	0.00	1603
mid	0.53	0.31	0.39	28483
high	0.63	0.83	0.72	42828
accuracy			0.61	72914
macro avg	0.72	0.38	0.37	72914
weighted avg	0.60	0.61	0.57	72914

Accuracy Score: 0.6084153934772472

Balanced Accuracy Score: 0.3799475125685053



Trained MLPClassifier in 630.3555474281311s

[19]: clf

```
[19]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer(strategy='median')),
                                                                              ('scaler',
                                                                              StandardScaler()))]),
                                                         ['age',
                                                                              'year_of_publication']),
                                                         ('cat',
                                                         OneHotEncoder(handle_unknown='ignore'),
                                                         ['book_title', 'book_author',
                                                                              'publisher', 'country'])])),
                        ('classifier',
```

```
MLPClassifier(alpha=0.001, max_iter=2, verbose=1))])
```

## 1.9 RandomForest Classifier

```
[23]: rfc = RandomForestClassifier(max_depth=50, verbose=1, n_jobs=-1)
```

```
[24]: clf = fit_model(algorithm=rfc, data=(X_train, X_test, y_train, y_test),  
    ↳preprocessor=preprocessor)
```

Started Training RandomForestClassifier on X\_train: (291656, 6) y\_train:  
(291656,)

```
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 2 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 46 tasks      | elapsed: 4.1min  
[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 8.8min finished
```

Evaluating model on X\_test: (72914, 6) y\_test: (72914,)

```
[Parallel(n_jobs=2)]: Using backend ThreadingBackend with 2 concurrent workers.  
[Parallel(n_jobs=2)]: Done 46 tasks      | elapsed: 0.6s  
[Parallel(n_jobs=2)]: Done 100 out of 100 | elapsed: 1.2s finished
```

Classification Report for RandomForestClassifier

```
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/_classification.py:1245:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/_classification.py:1245:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```


```
/usr/local/lib/python3.6/dist-packages/sklearn/metrics/_classification.py:1245:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

	precision	recall	f1-score	support
low	0.00	0.00	0.00	1603
mid	0.64	0.00	0.00	28483
high	0.59	1.00	0.74	42828

accuracy			0.59	72914
macro avg	0.41	0.33	0.25	72914
weighted avg	0.60	0.59	0.44	72914

Accuracy Score: 0.5876786351043696  
 Balanced Accuracy Score: 0.33361039593376624



Trained RandomForestClassifier in 537.0094470977783s

```
[25]: clf
```

```
[25]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer(strategy='median')),
                                                                              ('scaler',
                                                                              StandardScaler()))]),
                                                         ['age',
```

```

        'year_of_publication']],
        ('cat',
         OneHotEncoder(handle_unknown='ignore'),
         [
             ('book_title', 'book_author',
              'publisher', 'country')]
         ])),
        ('classifier',
         RandomForestClassifier(max_depth=50, n_jobs=-1, verbose=1))
    ])

```

## 1.10 Comparing Scalers

- StandardScaler
- MinMaxScaler
- RobustScaler
- MaxAbsScaler
- PowerTransformer
- QuantileTransformer
- Normalizer

```

[43]: from time import time

def fit_model_unscaled(algorithm, data):
    categorical_transformer = OneHotEncoder(handle_unknown='ignore')

    preprocessor = ColumnTransformer(
        transformers=[
            ('cat', categorical_transformer, categorical_features)
        ]
    )

    t1 = time()

    X_train, X_test, y_train, y_test = data

    clf = Pipeline(steps=[('preprocessor', preprocessor),
                           ('classifier', algorithm)])

    # train the model
    clf.fit(X_train, y_train)

    # test the model
    y_true = y_test.copy()
    y_pred = clf.predict(X_test)

    # get the classification report
    print(f"\nClassification Report for {algorithm.__class__.__name__}
    ↳Unscaled")

```

```

    print(classification_report(y_true, y_pred, target_names=target_names,
→labels=target_names))

    acc_score = accuracy_score(y_true, y_pred)
    bal_score = balanced_accuracy_score(y_true, y_pred)

    print(f"\nAccuracy Score: {acc_score}")
    print(f"Balanced Accuracy Score: {bal_score}")

    print()

    t2 = time()

    print(f'Trained {algorithm.__class__.__name__} in {(t2 - t1)}s')

    return clf

```

```

[25]: from time import time

def fit_model_scaler(algorithm, data, scaler):

    numeric_transformer = Pipeline(
        steps=[
            ('imputer', SimpleImputer(strategy='median')),
            ('scaler', scaler)]
    )

    categorical_transformer = OneHotEncoder(handle_unknown='ignore')

    preprocessor = ColumnTransformer(
        transformers=[
            ('num', numeric_transformer, numeric_features),
            ('cat', categorical_transformer, categorical_features)
        ]
    )

    t1 = time()

    X_train, X_test, y_train, y_test = data

    clf = Pipeline(steps=[('preprocessor', preprocessor),
                          ('classifier', algorithm)])

    # train the model
    clf.fit(X_train, y_train)

```



```

# test the model
y_true = y_test.copy()
y_pred = clf.predict(X_test)

# get the classification report
print(f"\nClassification Report for {algorithm.__class__.__name__} with
→{scaler.__class__.__name__}")
print(classification_report(y_true, y_pred, target_names=target_names,
→labels=target_names))

acc_score = accuracy_score(y_true, y_pred)
bal_score = balanced_accuracy_score(y_true, y_pred)

print(f"\nAccuracy Score: {acc_score}")
print(f"Balanced Accuracy Score: {bal_score}")

print()

t2 = time()

print(f'Trained {algorithm.__class__.__name__} in {(t2 - t1)}s')

return clf

```

```
[26]: algo = PassiveAggressiveClassifier(early_stopping=True, n_jobs=-1)
```

### 1.10.1 Standard Scaler

```
[27]: std_scaler = StandardScaler()
```

```
[28]: fit_model_scaler(algo, data=(X_train, X_test, y_train, y_test),
→scaler=std_scaler)
```

Classification Report for PassiveAggressiveClassifier with StandardScaler

	precision	recall	f1-score	support
low	0.06	0.06	0.06	1603
mid	0.46	0.36	0.40	28483
high	0.62	0.72	0.67	42828
accuracy			0.56	72914
macro avg	0.38	0.38	0.38	72914
weighted avg	0.55	0.56	0.55	72914

Accuracy Score: 0.5615656801163014

Balanced Accuracy Score: 0.3785619997072959

Trained PassiveAggressiveClassifier in 7.8669610023498535s

```
[28]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer(strategy='median')),
                                                                              StandardScaler())])),
                        ('cat',
                        OneHotEncoder(handle_unknown='ignore'),
                        ('classifier',
                        PassiveAggressiveClassifier(early_stopping=True, n_jobs=-1)))]),
                        [('age',
                          'year_of_publication']],
                        [('book_title', 'book_author',
                          'publisher', 'country')]))],
```

### 1.10.2 MinMax Scaler

```
[29]: minmax_scaler = MinMaxScaler()
```

```
[30]: fit_model_scaler(algo, data=(X_train, X_test, y_train, y_test),
    ↪ scaler=minmax_scaler)
```

Classification Report for PassiveAggressiveClassifier with MinMaxScaler

	precision	recall	f1-score	support
low	0.05	0.00	0.01	1603
mid	0.44	0.55	0.49	28483
high	0.64	0.55	0.59	42828
accuracy			0.54	72914
macro avg	0.37	0.37	0.36	72914
weighted avg	0.55	0.54	0.54	72914

Accuracy Score: 0.5409249252544093

Balanced Accuracy Score: 0.3694473636392875

Trained PassiveAggressiveClassifier in 7.548543214797974s

```
[30]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer(strategy='median')),
```

```

('scaler',
MinMaxScaler()))],
['age',
'year_of_publication']],
('cat',
OneHotEncoder(handle_unknown='ignore'),
['book_title', 'book_author',
'publisher', 'country']]])),
('classifier',
PassiveAggressiveClassifier(early_stopping=True, n_jobs=-1))]

```

### 1.10.3 MaxAbs Scaler

```

[31]: maxabs_scaler = MaxAbsScaler()
[32]: fit_model_scaler(algo, data=(X_train, X_test, y_train, y_test),
    →scaler=maxabs_scaler)

```

Classification Report for PassiveAggressiveClassifier with MaxAbsScaler

	precision	recall	f1-score	support
low	0.07	0.01	0.02	1603
mid	0.49	0.19	0.27	28483
high	0.61	0.87	0.72	42828
accuracy			0.59	72914
macro avg	0.39	0.36	0.34	72914
weighted avg	0.55	0.59	0.53	72914

Accuracy Score: 0.5882820857448501

Balanced Accuracy Score: 0.3598257153815126

Trained PassiveAggressiveClassifier in 7.605406761169434s

```

[32]: Pipeline(steps=[('preprocessor',
    ColumnTransformer(transformers=[('num',
    Pipeline(steps=[('imputer',
SimpleImputer(strategy='median')),
('scaler',
MaxAbsScaler()))],
['age',
'year_of_publication']],
('cat',
OneHotEncoder(handle_unknown='ignore'),
['book_title', 'book_author',

```

```

('publisher', 'country']]])),
('classifier',
 PassiveAggressiveClassifier(early_stopping=True, n_jobs=-1)))

```

#### 1.10.4 Robust Scaler

```
[33]: robust_scaler = RobustScaler(quantile_range=(25, 75))
```

```
[34]: fit_model_scaler(algo, data=(X_train, X_test, y_train, y_test),
    ↪ scaler=robust_scaler)
```

Classification Report for PassiveAggressiveClassifier with RobustScaler

	precision	recall	f1-score	support
low	0.08	0.03	0.05	1603
mid	0.47	0.26	0.33	28483
high	0.61	0.81	0.70	42828
accuracy			0.58	72914
macro avg	0.39	0.37	0.36	72914
weighted avg	0.54	0.58	0.54	72914

Accuracy Score: 0.5769811010231232

Balanced Accuracy Score: 0.36651556542661007

Trained PassiveAggressiveClassifier in 7.649526119232178s

```
[34]: Pipeline(steps=[('preprocessor',
    ColumnTransformer(transformers=[('num',
    Pipeline(steps=[('imputer',
    SimpleImputer(strategy='median')),
    RobustScaler(quantile_range=(25,
    75)))]),
    ('scaler',
    ['age',
    'year_of_publication']),
    ('cat',
    OneHotEncoder(handle_unknown='ignore'),
    ['book_title', 'book_author',
    'publisher', 'country']]])),
('classifier',
 PassiveAggressiveClassifier(early_stopping=True, n_jobs=-1)))

```

### 1.10.5 PowerTransform (Yeo-Johnson)

```
[35]: power_trans = PowerTransformer(method='yeo-johnson')
```

```
[36]: fit_model_scaler(algo, data=(X_train, X_test, y_train, y_test),  
    ↪ scaler=power_trans)
```

```
/usr/local/lib/python3.6/dist-packages/numpy/core/_methods.py:205:
```

```
RuntimeWarning: overflow encountered in multiply
```

```
  x = um.multiply(x, x, out=x)
```

```
/usr/local/lib/python3.6/dist-packages/numpy/core/_methods.py:216:
```

```
RuntimeWarning: overflow encountered in reduce
```

```
  ret = umr_sum(x, axis, dtype, out, keepdims)
```

Classification Report for PassiveAggressiveClassifier with PowerTransformer

	precision	recall	f1-score	support
low	0.06	0.04	0.05	1603
mid	0.43	0.58	0.50	28483
high	0.64	0.51	0.57	42828
accuracy			0.53	72914
macro avg	0.38	0.37	0.37	72914
weighted avg	0.55	0.53	0.53	72914

Accuracy Score: 0.5252214938146309

Balanced Accuracy Score: 0.37488500156296717

Trained PassiveAggressiveClassifier in 10.008253574371338s

```
[36]: Pipeline(steps=[('preprocessor',  
    ColumnTransformer(transformers=[('num',  
    Pipeline(steps=[('imputer',  
    SimpleImputer(strategy='median')),  
    ('scaler',  
    PowerTransformer()))]),  
    [ 'age',  
    'year_of_publication']),  
    ('cat',  
    OneHotEncoder(handle_unknown='ignore'),  
    [ 'book_title', 'book_author',  
    'publisher', 'country'])])),  
    ('classifier',  
    PassiveAggressiveClassifier(early_stopping=True, n_jobs=-1)))]
```

### 1.10.6 Quantile Transform

```
[37]: quant_trans = QuantileTransformer(output_distribution='normal')
```

```
[38]: fit_model_scaler(algo, data=(X_train, X_test, y_train, y_test),  
    ↪ scaler=quant_trans)
```

Classification Report for PassiveAggressiveClassifier with QuantileTransformer

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

low	0.08	0.03	0.05	1603
mid	0.43	0.59	0.50	28483
high	0.64	0.49	0.56	42828
accuracy			0.52	72914
macro avg	0.38	0.37	0.37	72914
weighted avg	0.54	0.52	0.52	72914

Accuracy Score: 0.5197492936884549

Balanced Accuracy Score: 0.37149925184802246

Trained PassiveAggressiveClassifier in 8.048646688461304s

```
[38]: Pipeline(steps=[('preprocessor',  
    ColumnTransformer(transformers=[('num',  
    Pipeline(steps=[('imputer',  
    SimpleImputer(strategy='median')),  
    ('scaler',  
    QuantileTransformer(output_distribution='normal'))]),  
    ['age',  
    'year_of_publication']),  
    ('cat',  
    OneHotEncoder(handle_unknown='ignore'),  
    ['book_title', 'book_author',  
    'publisher', 'country'])])),  
    ('classifier',  
    PassiveAggressiveClassifier(early_stopping=True, n_jobs=-1))])
```

### 1.10.7 Normalizer (L2)

```
[39]: norm_scaler = Normalizer()
```

```
[40]: fit_model_scaler(algo, data=(X_train, X_test, y_train, y_test),  
    ↪ scaler=norm_scaler)
```

### Classification Report for PassiveAggressiveClassifier with Normalizer

	precision	recall	f1-score	support
low	0.05	0.03	0.04	1603
mid	0.46	0.30	0.36	28483
high	0.62	0.77	0.69	42828
accuracy			0.57	72914
macro avg	0.38	0.37	0.36	72914
weighted avg	0.54	0.57	0.55	72914

Accuracy Score: 0.5706722988726445

Balanced Accuracy Score: 0.36646952962846085

Trained PassiveAggressiveClassifier in 8.04029631614685s

```
[40]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('imputer',
                                                                              SimpleImputer(strategy='median')),
                                                                              ('scaler',
                                                                              Normalizer()))]),
                                                         ['age',
                                                                              'year_of_publication']),
                                                         ('cat',
                                                         OneHotEncoder(handle_unknown='ignore'),
                                                         ['book_title', 'book_author',
                                                                              'publisher', 'country'])])),
                        ('classifier',
                        PassiveAggressiveClassifier(early_stopping=True, n_jobs=-1))])
```

### 1.10.8 Unscaled

```
[44]: fit_model_unscaled(algo, data=(X_train, X_test, y_train, y_test))
```

### Classification Report for PassiveAggressiveClassifier Unscaled

	precision	recall	f1-score	support
low	0.09	0.02	0.03	1603
mid	0.48	0.24	0.32	28483
high	0.61	0.83	0.70	42828
accuracy			0.58	72914
macro avg	0.39	0.36	0.35	72914

weighted avg      0.55      0.58      0.54      72914

Accuracy Score: 0.5831938996626163

Balanced Accuracy Score: 0.36462720452376884

Trained PassiveAggressiveClassifier in 7.465895652770996s

```
[44]: Pipeline(steps=[('preprocessor',  
                        ColumnTransformer(transformers=[('cat',  
                                                         OneHotEncoder(handle_unknown='ignore'),  
                                                         ['book_title', 'book_author',  
                                                         'publisher', 'country'])])),  
                  ('classifier',  
                   PassiveAggressiveClassifier(early_stopping=True, n_jobs=-1))])
```

```
[ ]:
```



# BookReview-Classification-Unsupervised

January 20, 2021

```
[!]: ! pip install --upgrade scikit-learn
[!]: ! pip install scikit-learn-extra
[!]: ! pip install hdbscan
[!]: ! python -m spacy download en_core_web_sm
```

---

## 1 Book Crossing - Classification (Unsupervised)

```
[3]: %matplotlib inline

import scipy
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

sns.set()
palette = sns.color_palette("icefire")

plt.style.use('ggplot')

sns.set_context("talk")

[127]: dataset = pd.read_csv('book_crossing.classification.cleaned.csv')

[128]: dataset['age'] = dataset['age'].astype(np.float64)
dataset['book_rating'] = dataset['book_rating'].astype('category')
dataset['book_title'] = dataset['book_title'].astype('category')
dataset['book_author'] = dataset['book_author'].astype('category')
dataset['year_of_publication'] = dataset['year_of_publication'].astype(np.
    →float64)
dataset['publisher'] = dataset['publisher'].astype('category')
dataset['country'] = dataset['country'].astype('category')
```

```
[136]: dataset['book_title'].cat.categories.shape
```

```
[136]: (132033,)
```

```
[139]: dataset['book_author'].cat.categories.shape
```

```
[139]: (60652,)
```

```
[140]: dataset['publisher'].cat.categories.shape
```

```
[140]: (11311,)
```

```
[141]: dataset['country'].cat.categories.shape
```

```
[141]: (51,)
```

```
[6]: dataset.head()
```

```
[6]:   age book_rating ... publisher country
0  34.0         mid ... HarperFlamingo Canada  canada
1  30.0         high ... HarperFlamingo Canada  canada
2  34.0         high ... HarperFlamingo Canada  canada
3  34.0         high ... HarperFlamingo Canada  canada
4  34.0         high ... HarperFlamingo Canada  canada
```

```
[5 rows x 7 columns]
```

```
[7]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 364570 entries, 0 to 364569
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	age	364570 non-null	float64
1	book_rating	364570 non-null	category
2	book_title	364570 non-null	category
3	book_author	364570 non-null	category
4	year_of_publication	364570 non-null	float64
5	publisher	364570 non-null	category
6	country	364570 non-null	category

```
dtypes: category(5), float64(2)
```

```
memory usage: 19.1 MB
```

```
[8]: c_dataset = dataset["book_title"].astype(str) + " " + \
      dataset["book_author"].astype(str) + " " + \
      dataset["publisher"].astype(str) + " " + \
      dataset["year_of_publication"].astype(str) + " " + \
      dataset["age"].astype(str) + " " + \
      dataset["country"].astype(str)
```

```
[9]: for ex in c_dataset.sample(frac=0.2)[:5]:
      print(ex)
```

The Vanishing Vampire (The Accidental Monsters , No 1) David Lubar Scholastic  
 1997.0 12.0 usa  
 REMEMBER ME Mary Higgins Clark Simon & Schuster 1994.0 34.0 usa  
 The Mummy or Ramses the Damned Anne Rice Ballantine Books 1991.0 22.0 usa  
 Bittersweet Rain Sandra Brown Warner Books 2000.0 34.0 usa  
 Arthur Stephen R. Lawhead Zondervan Publishing Company 1996.0 21.0 usa

```
[92]: small_dataset = dataset.copy().sample(frac=0.03)
```

```
[93]: small_dataset.shape
```

```
[93]: (10937, 7)
```

```
[94]: c_dataset_small = small_dataset["book_title"].astype(str) + " " + \
      small_dataset["book_author"].astype(str) + " " + \
      small_dataset["publisher"].astype(str) + " " + \
      small_dataset["year_of_publication"].astype(str) + " " + \
      small_dataset["age"].astype(str) + " " + \
      small_dataset["country"].astype(str)
```

```
[95]: import spacy
      from sklearn.base import BaseEstimator, TransformerMixin
      import en_core_web_sm

      nlp = en_core_web_sm.load()
```

```
[96]: from sklearn import set_config
      from sklearn.compose import ColumnTransformer
      from sklearn.pipeline import Pipeline
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import classification_report, plot_confusion_matrix, \
      →confusion_matrix, accuracy_score, balanced_accuracy_score
      from sklearn.cluster import KMeans, DBSCAN, Birch, MiniBatchKMeans, \
      →SpectralClustering, AgglomerativeClustering, MeanShift, AffinityPropagation, \
      →OPTICS
      from sklearn.decomposition import TruncatedSVD

      set_config(display='diagram')
```

```
[97]: class SpacyVectorTransformer(BaseEstimator, TransformerMixin):
      def __init__(self, nlp):
          self.nlp = nlp
          self.dim = 300

      def fit(self, X, y):
          return self

      def transform(self, X):
          # Doc.vector defaults to an average of the token vectors.
          # https://spacy.io/api/doc#vector
```

```
return [self.nlp(text).vector for text in X]
```

```
[98]: X, y = c_dataset, dataset['book_rating']
```

```
[99]: X_small, y_small = c_dataset_small, small_dataset['book_rating']
```

```
[100]: target_names = ['low', 'mid', 'high']
```

## 1.1 Unsupervised Models

- KMedoids
- AgglomerativeClustering
- DBSCAN
- KMeansClustering
- HDBSCAN
- MiniBatchKMeans

```
[102]: """  
km = KMeans(n_clusters=3) # works (needs remap of output)  
dbscan = DBSCAN()  
birch = Birch() # crash  
mbkm = MiniBatchKMeans() # works (needs remap of output)  
sc = SpectralClustering() # crash  
ac = AgglomerativeClustering() # sparse not supported  
ms = MeanShift(bandwidth=3) # sparse not supported  
ap = AffinityPropagation() # crash  
oo = OPTICS() # sparse not supported  
"""
```

```
[102]: '\nkm = KMeans(n_clusters=3) # works (needs remap of output)\nndbscan =  
DBSCAN()\nbirch = Birch() # crash\nmbkm = MiniBatchKMeans() # works (needs remap  
of output)\nsc = SpectralClustering() # crash\nnac = AgglomerativeClustering() #  
sparse not supported\nms = MeanShift(bandwidth=3) # sparse not supported\nnap =  
AffinityPropagation() # crash\nnoo = OPTICS() # sparse not supported\n'
```

```
[103]: from time import time  
  
def fit_model(algorithm, data):  
  
    t1 = time()  
  
    X, y = data  
  
    print(f'\nStarted Training {algorithm.__class__.__name__} on X: {X.shape} y:  
→ {y.shape}')  
  
    embeddings_pipeline = Pipeline(  
        steps=[  
            ("mean_embeddings", SpacyVectorTransformer(nlp=nlp)),
```

```

        ("reduce_dim", TruncatedSVD(50)),
        ("clusterer", algorithm),
    ]
)

# train the model
embeddings_pipeline.fit(X, y)

print(f"\nEvaluating model on X_test: {X.shape} y_test: {y.shape}")

# test the model
y_true = y.copy()

if isinstance(embeddings_pipeline['clusterer'], (AgglomerativeClustering,
→DBSCAN, OPTICS, HDBSCAN)):
    y_pred = embeddings_pipeline['clusterer'].labels_
else:
    y_pred = embeddings_pipeline.predict(X)

y_pred = np.array(list(map(lambda x: "low" if x == 0 else "mid" if x == 1
→else "high", y_pred)))

# get the classification report
print(f"\nClassification Report for {algorithm.__class__.__name__}")
print(classification_report(y_true, y_pred, target_names=target_names,
→labels=target_names))

acc_score = accuracy_score(y_true, y_pred)
bal_score = balanced_accuracy_score(y_true, y_pred)

print(f"\nAccuracy Score: {acc_score}")
print(f"Balanced Accuracy Score: {bal_score}")

print()
# show the confusion matrix
cmmat_table = pd.DataFrame({'y_true': y_true, 'y_pred': y_pred})
conmat = pd.crosstab(cmmat_table.y_true, cmmat_table.y_pred,
→rownames=['Actual'], colnames=['Predicted'], margins=True, normalize='all')
ax = plt.axes()
sns.set(rc={'figure.figsize': (9, 7)})
sns.heatmap(conmat, annot=True, ax=ax)
ax.set_title(f'{algorithm.__class__.__name__}')
plt.show()
print()

t2 = time()

```

```

print(f'Trained {algorithm.__class__.__name__} in {(t2 - t1)}s')

return embeddings_pipeline

```

## 1.2 KMedoids Clustering

```

[107]: from sklearn_extra.cluster import KMedoids
[108]: kmedoids = KMedoids(n_clusters=3, max_iter=1)
[109]: clf = fit_model(algorithm=kmedoids, data=(X_small, y_small))

```

Started Training KMedoids on X: (10937,) y: (10937,)


Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for KMedoids

	precision	recall	f1-score	support
low	0.02	0.34	0.04	231
mid	0.38	0.30	0.34	4333
high	0.60	0.31	0.41	6373
accuracy			0.31	10937
macro avg	0.33	0.32	0.26	10937
weighted avg	0.50	0.31	0.37	10937

Accuracy Score: 0.30785407332906645

Balanced Accuracy Score: 0.3184037656596693



Trained KMedoids in 223.7921495437622s

```
[110]: clf
```

```
[110]: Pipeline(steps=[('mean_embeddings',
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at
0x7ff253d3b6d8>)),
                      ('reduce_dim', TruncatedSVD(n_components=50)),
                      ('clusterer', KMedoids(max_iter=1, n_clusters=3))])
```

### 1.3 Agglomerative Clustering

```
[111]: aggcl = AgglomerativeClustering(n_clusters=3)
```

```
[112]: clf = fit_model(algorithm=aggcl, data=(X_small, y_small))
```

Started Training AgglomerativeClustering on X: (10937,) y: (10937,)


Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for AgglomerativeClustering

	precision	recall	f1-score	support
low	0.02	0.37	0.04	231
mid	0.37	0.31	0.34	4333
high	0.57	0.24	0.34	6373
accuracy			0.27	10937
macro avg	0.32	0.31	0.24	10937
weighted avg	0.48	0.27	0.33	10937

Accuracy Score: 0.2729267623662796

Balanced Accuracy Score: 0.30877414843520845





Trained AgglomerativeClustering in 119.64335346221924s

```
[113]: clf
```

```
[113]: Pipeline(steps=[('mean_embeddings',  
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at  
                        0x7ff253d3b6d8>)),  
                      ('reduce_dim', TruncatedSVD(n_components=50)),  
                      ('clusterer', AgglomerativeClustering(n_clusters=3))])
```

## 1.4 DBSCAN

```
[114]: dbscan = DBSCAN(n_jobs=-1)
```

```
[115]: clf = fit_model(algorithm=dbscan, data=(X_small, y_small))
```

Started Training DBSCAN on X: (10937,) y: (10937,)


Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for DBSCAN

	precision	recall	f1-score	support
low	0.00	0.00	0.00	231
mid	0.20	0.00	0.00	4333
high	0.58	1.00	0.74	6373
accuracy			0.58	10937
macro avg	0.26	0.33	0.25	10937
weighted avg	0.42	0.58	0.43	10937

Accuracy Score: 0.5818780287098839

Balanced Accuracy Score: 0.3328872224199264



Trained DBSCAN in 116.67007970809937s

```
[116]: clf
```

```
[116]: Pipeline(steps=[('mean_embeddings',
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at
0x7ff253d3b6d8>)),
                      ('reduce_dim', TruncatedSVD(n_components=50)),
                      ('clusterer', DBSCAN(n_jobs=-1))])
```

## 1.5 KMeans Clustering

```
[117]: kmeans = KMeans(n_clusters=3, n_init=3)
```

```
[118]: clf = fit_model(algorithm=kmeans, data=(X_small, y_small))
```

Started Training KMeans on X: (10937,) y: (10937,)


Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for KMeans

	precision	recall	f1-score	support
low	0.02	0.40	0.04	231
mid	0.39	0.26	0.32	4333
high	0.60	0.35	0.44	6373
accuracy			0.32	10937
macro avg	0.34	0.34	0.27	10937
weighted avg	0.50	0.32	0.38	10937

Accuracy Score: 0.31535155892840816

Balanced Accuracy Score: 0.33658550914710594



Trained KMeans in 222.76584196090698s

```
[119]: clf
```

```
[119]: Pipeline(steps=[('mean_embeddings',  
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at  
                        0x7ff253d3b6d8>)),  
                        ('reduce_dim', TruncatedSVD(n_components=50)),  
                        ('clusterer', KMeans(n_clusters=3, n_init=3))])
```

## 1.6 HDBSCAN

```
[120]: from hdbscan import HDBSCAN
```

```
[121]: hscan = HDBSCAN(min_cluster_size=3)
```

```
[122]: clf = fit_model(algorithm=hscan, data=(X_small, y_small))
```

Started Training HDBSCAN on X: (10937,) y: (10937,)

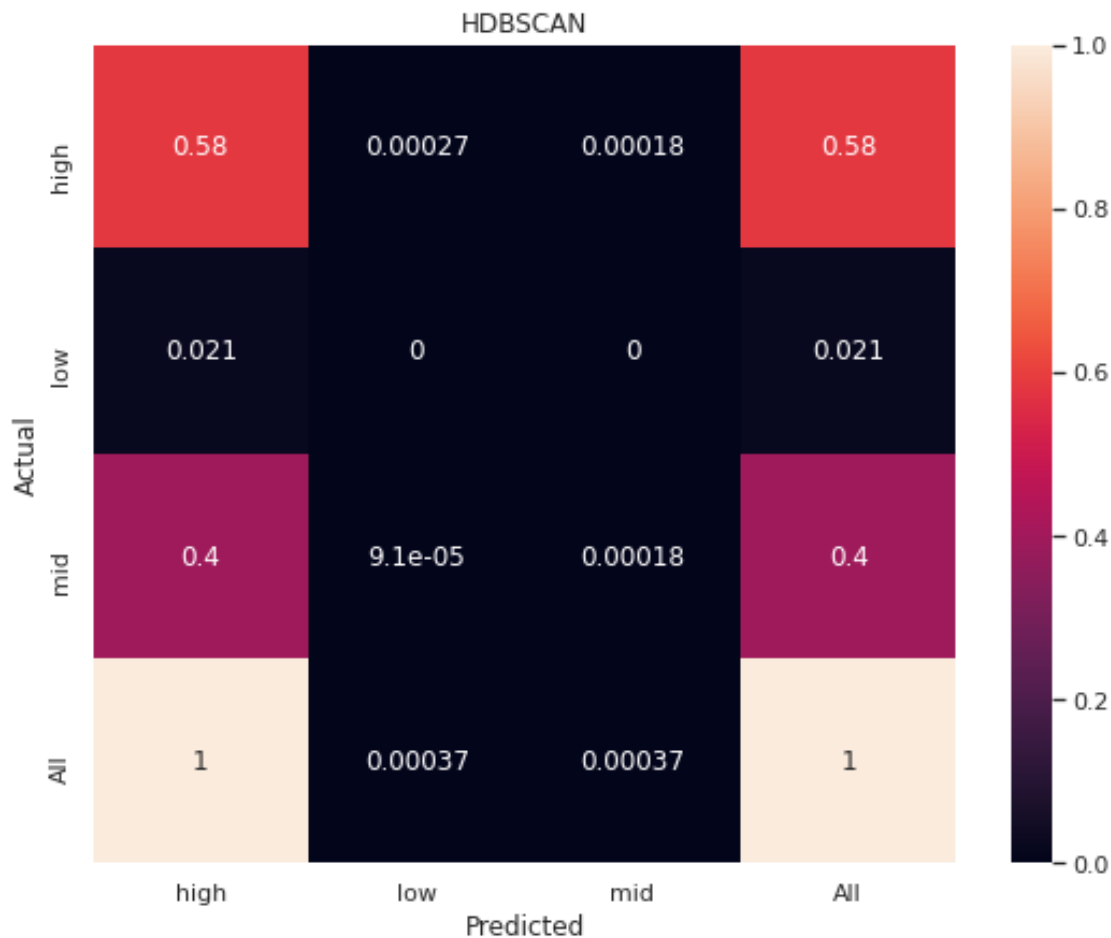
Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for HDBSCAN

	precision	recall	f1-score	support
low	0.00	0.00	0.00	231
mid	0.50	0.00	0.00	4333
high	0.58	1.00	0.74	6373
accuracy			0.58	10937
macro avg	0.36	0.33	0.25	10937
weighted avg	0.54	0.58	0.43	10937

Accuracy Score: 0.5824266252171528

Balanced Accuracy Score: 0.3332256713684369



Trained HDBSCAN in 121.39391756057739s

```
[123]: clf
```

```
[123]: Pipeline(steps=[('mean_embeddings',
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at
0x7ff253d3b6d8>)),
                        ('reduce_dim', TruncatedSVD(n_components=50)),
                        ('clusterer', HDBSCAN(min_cluster_size=3))])
```

## 1.7 Mini Batch KMeans

```
[124]: mb_kmeans = MiniBatchKMeans(n_clusters=3, max_iter=1)
```

```
[125]: clf = fit_model(algorithm=mb_kmeans, data=(X_small, y_small))
```

Started Training MiniBatchKMeans on X: (10937,) y: (10937,)

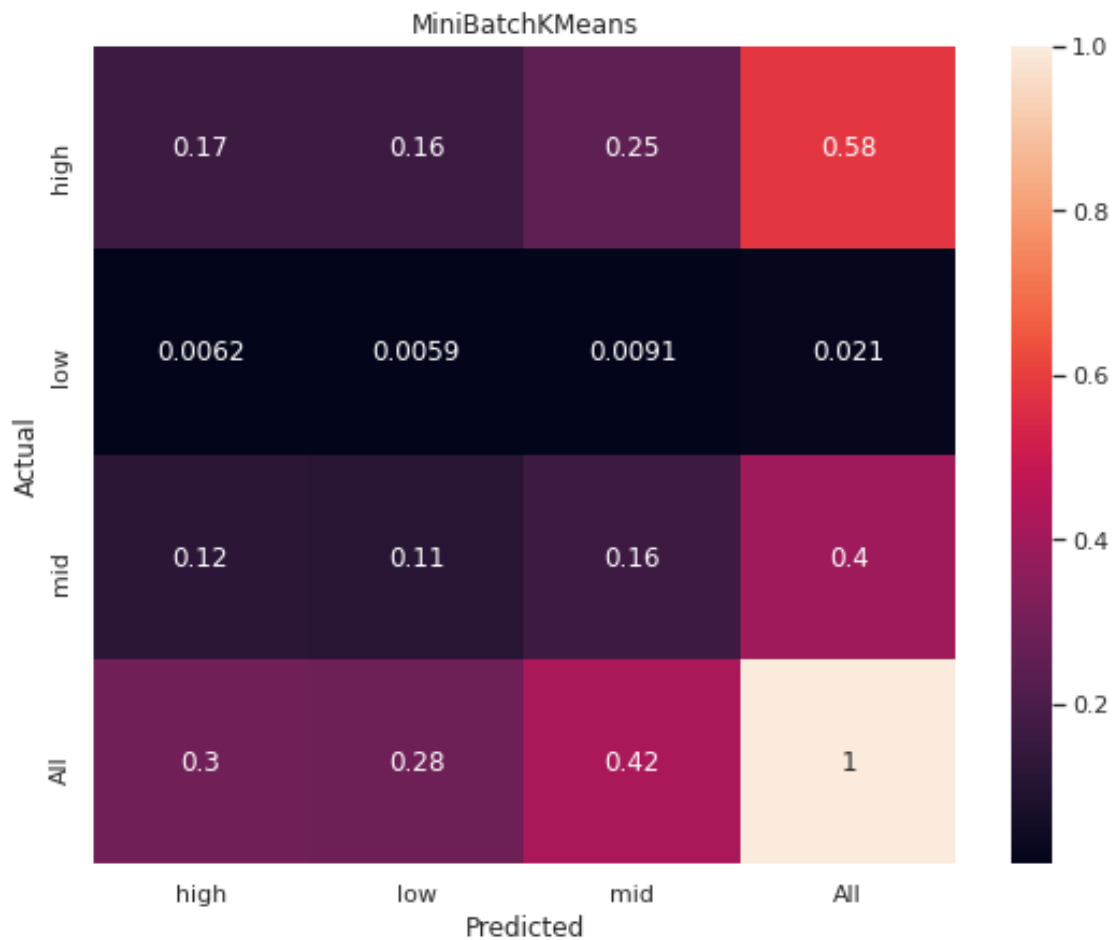
Evaluating model on X\_test: (10937,) y\_test: (10937,)

Classification Report for MiniBatchKMeans

	precision	recall	f1-score	support
low	0.02	0.28	0.04	231
mid	0.39	0.41	0.40	4333
high	0.57	0.29	0.38	6373
accuracy			0.34	10937
macro avg	0.33	0.33	0.27	10937
weighted avg	0.49	0.34	0.38	10937

Accuracy Score: 0.3378440157264332

Balanced Accuracy Score: 0.32632201462411303



Trained MiniBatchKMeans in 217.56095910072327s

```
[126]: clf
```

```
[126]: Pipeline(steps=[('mean_embeddings',  
                        SpacyVectorTransformer(nlp=<spacy.lang.en.English object at  
0x7ff253d3b6d8>)),  
                      ('reduce_dim', TruncatedSVD(n_components=50)),  
                      ('clusterer', MiniBatchKMeans(max_iter=1, n_clusters=3))])
```

```
[ ]:
```